Protecting Victims of Authorised Push Payment Fraud

1 Preliminaries

1.1 Notations and Assumptions

We use $\texttt{Enc}(.)$ and $\texttt{Dec}(.)$ to denote the encrypting and decrypting algorithms of a semantically secure symmetric key encryption scheme. In this work, we use a committee of honest arbiters $\mathcal{D} : \{\mathcal{D}_1, ..., \mathcal{D}_n\}$. Each arbiter, given a set of inputs, provides a binary verdict. We assume $\mathcal{D}_i$'s share a pair of keys $(pk_{\mathcal{D}}, sk_{\mathcal{D}})$ of an asymmetric key encryption. The encryption scheme has key generating $\tilde{\texttt{keyGen}}(1^\lambda) \rightarrow (sk_{\mathcal{D}}, pk_{\mathcal{D}})$, encrypting $\tilde{\texttt{Enc}}(pk_{\mathcal{D}}, .)$ and decrypting $\tilde{\texttt{Dec}}(sk_{\mathcal{D}}, .)$ algorithms, where its public key is known to everyone. Also, we assume all arbiters have interacted with each other to agree on a secret key, $\bar{k}_0$. We use ϕ to denote a null value. Our proposed solution is built upon the existing online banking system. We assume the banking system has three algorithms in place, (1) $\texttt{setupNewPayee}(.)$: allows a customer to add a new payee to its payee list, (2) $\texttt{ammendExistingPayee}(.)$: lets a customer amend its existing payee list, and (3) $\texttt{pay}(.)$: transfers money from the customer's account to a payee's account that is specific by the customer. We denote the inputs of algorithms $\texttt{setupNewPayee}(.), \texttt{ammendExistingPayee}(.)$, and $\texttt{pay}(.)$ by in_s, in_a, and in_p respectively. We assume the source code of the online banking system is static, is not updated, and any change to the source code can be detected, e.g., the bank can use a cryptographic commitment to commit to the source code. Moreover, we assume the online banking system is secure.

1.2 Digital Signature

A digital signature is a scheme for verifying the authenticity of digital messages or documents. Below, we restate its formal definition, taken from [7].

Definition 1. *A signature scheme involves three sub-processes,* $\texttt{Signature} := (\texttt{Sig.keyGen}, \texttt{Sig.sign}, \texttt{Sig.ver})$*, that are defined as follows.*

- $\texttt{Sig.keyGen}(1^\lambda) \rightarrow (sk, pk)$*. A probabilistic algorithm run by a signer. It takes as input a security parameter. It outputs a parameter pair:* (sk, pk)*, consisting of secret:* sk*, and public:* pk *parameters.*
- $\texttt{Sig.sign}(sk, pk, u) \rightarrow sig$*. An algorithm run the signer. It takes as input parameter pair:* (sk, pk) *and a file:* u*. It outputs a signature:* sig*.*
- $\texttt{Sig.ver}(pk, u, sig) \rightarrow h \in \{0, 1\}$*. A deterministic algorithm run by a verifier. It takes as input public parameter:* pk*, file:* u*, and signature:* sig*. It checks the signature's validity. If the verification passes, then it outputs* 1*; otherwise, it outputs* 0*.*

A digital signature scheme should meet the following properties:

- *Correctness.* For every input u it holds that:

$$Pr\Big[\ \texttt{Sig.ver}(pk, u, \texttt{Sig.sign}(sk, pk, u)) = 1\ :\ \texttt{Sig.keyGen}(1^\lambda) \rightarrow (sk, pk)\Big] = 1$$

- *Existential unforgeability under chosen message attacks.* A (PPT) adversary that obtains pk and has access to a signing oracle for messages of its choice, cannot create a valid pair (u^*, sig^*) for a new file u^* (that was never a query to the signing oracle), except with a small probability, σ. More formally:

$$\Pr\begin{bmatrix} u^* \notin Q \wedge & \texttt{Cer.keyGen}(1^\lambda) \rightarrow (sk, pk) \\ \texttt{Sig.ver}(pk, u^*, sig^*) = 1 & : \mathcal{A}^{\texttt{Sig.sign}(k,)}(pk) \rightarrow (u^*, sig^*) \end{bmatrix} \leq \mu(\lambda)$$

where Q is the set of queries that $\mathcal{A}$ sent to the certificate generator oracle.

An application of a digital signature is in digital certificate, which is a document digitally signed by a certificate generator. Given a certificate and its parameters, anyone can check whether it has been correctly generated by a valid generator. There is a case where a *hard copy* certificate is used. In this case, the process $\texttt{Sig.keyGen}(.)$ outputs a blank legitimate stamped certificate as a private parameter: sk, and the description of a standard legitimate certificate as a public parameter: pk. Moreover, the process $\texttt{Sig.sign}(k, u)$ takes k and the file u on which a certificate should be generated and outputs a stamped certificate with the information printed on it. The process $\texttt{Sig.ver}(pk, u, sig)$ takes the public parameter, the file, and the hard copy of the certificate and outputs 1 if it is valid and 0 if it is not. Note, when a hard copy certificate is considered, it is not possible to precisely define the success probability of the adversary, as it depends on the technology available to the adversary to generate a blank stamped certificate that looks like a legitimate one. In the real world however, this probability is usually small (but it may not be negligible). In this paper, we mainly consider a digital certificate; however, our solution can adopt hard copy certificates as well with the above caveat regarding the adversary's success probability.

1.3 Smart Contract

Cryptocurrencies, such as Bitcoin [10] and Ethereum [14], beyond offering a decentralised currency, support computations on transactions. In this setting, often a certain computation logic is encoded in a computer program, called a *"smart contract"*. Although Bitcoin, the first decentralised cryptocurrency, supports smart contracts, the functionality of Bitcoin's smart contracts is very limited, due to the use of the underlying programming language that does not support arbitrary tasks. To address this limitation, Ethereum, as a generic smart contract platform, was designed. Thus far, Ethereum has been the most predominant cryptocurrency framework that lets users define arbitrary smart. This framework allows users to create an account with a unique account number or address. Such users are often called external account holders, which can send (or deploy) their contracts to the framework's blockchain. In this framework, a contract's code and its related data is held by every node in the blockchain's network. Ethereum smart contracts are often written in a high-level Turing-complete programming language called "Solidity".The program execution's correctness is guaranteed by the security of the underlying blockchain components. To prevent a denial of service attack, the framework requires a transaction creator to pay a fee, called *"gas"*, depending on the complexity of the contract running on it.

1.4 Commitment Scheme

A commitment scheme involves two parties, *sender* and *receiver*, and includes two phases: *commit* and *open*. In the commit phase, the sender commits to a message: x as $\texttt{Com}(x, r) = \texttt{Com}_x$, that involves a secret value: $r \xleftarrow{\$} \{0,1\}^\lambda$. In the end of the commit phase, the commitment $\texttt{Com}_x$ is sent to the receiver. In the open phase, the sender sends the opening $\ddot{x} := (x, r)$ to the receiver who verifies its correctness: $\texttt{Ver}(\texttt{Com}_x, \ddot{x}) \stackrel{?}{=} 1$ and accepts if the output is 1. A commitment scheme must satisfy two properties: (a) *hiding*: it is infeasible for an adversary (i.e., the receiver) to learn any information about the committed message x, until the commitment $\texttt{Com}_x$ is opened, and (b) *binding*: it is infeasible for an adversary (i.e., the sender) to open a commitment $\texttt{Com}_x$ to different values $\ddot{x}' := (x', r')$ than that was used in the commit phase, i.e., infeasible to find $\ddot{x}'$, *s.t.* $\texttt{Ver}(\texttt{Com}_x, \ddot{x}) = \texttt{Ver}(\texttt{Com}_x, \ddot{x}') = 1$, where $\ddot{x} \neq \ddot{x}'$. There exist efficient non-interactive commitment schemes both in (a) the standard model, e.g., Pedersen scheme [11], and (b) the random oracle model using the well-known hash-based scheme such that committing is : $\texttt{H}(x||r) = \texttt{Com}_x$ and $\texttt{Ver}(\texttt{Com}_x, \ddot{x})$ requires checking: $\texttt{H}(x||r) \stackrel{?}{=} \texttt{Com}_x$, where $\texttt{H} : \{0,1\}^* \rightarrow \{0,1\}^\lambda$ is a collision resistant hash function; i.e., the probability to find x and x' such that $\texttt{H}(x) = \texttt{H}(x')$ is negligible in the security parameter, λ.

1.5 Statement Agreement Protocol

Recently, a scheme called "Statement Agreement Protocol" (SAP) has been proposed in [1]. It lets two mutually distrusted parties, e.g., $\mathcal{B}$ and $\mathcal{C}$, efficiently agree on a private statement, π. Informally, the SAP

satisfies the following four properties: (1) neither party can convince a third-party verifier that it has agreed with its counter-party on a different statement than the one both parties previously agreed on, (2) after they agree on a statement, an honest party can (almost) always prove to the verifier that it has the agreement, (3) the privacy of the statement is preserved (from the public), and (4) after both parties reach an agreement, neither can deny it. The SAP relies on a smart contract and commitment scheme and assumes that each party already has a blockchain public address, $adr_{\mathcal{R}}$ (where $\mathcal{R} \in \{\mathcal{B}, \mathcal{C}\}$). Below, we restate the SAP, taken from [1].

1. **Initiate**. $\texttt{SAP.init}(1^\lambda, adr_{\mathcal{B}}, adr_{\mathcal{C}}, \pi)$
   The following steps are taken by $\mathcal{B}$.
   (a) Deploys a smart contract that explicitly states both parties' addresses, $adr_{\mathcal{B}}$ and $adr_{\mathcal{C}}$. Let adr_{SAP} be the deployed contract's address.
   (b) Picks a random value r, and commits to the statement, $\texttt{Com}(\pi, r) = g_{\mathcal{B}}$.
   (c) Sends adr_{SAP} and $\ddot{\pi} := (\pi, r)$ to $\mathcal{C}$, and $g_{\mathcal{B}}$ to the contract.
2. **Agreement**. $\texttt{SAP.agree}(\pi, r, g_{\mathcal{B}}, adr_{\mathcal{B}}, adr_{\text{SAP}})$
   The following steps are taken by $\mathcal{C}$.
   (a) Checks if $g_{\mathcal{B}}$ was sent from $adr_{\mathcal{B}}$, and checks locally $\texttt{Ver}(g_{\mathcal{B}}, \ddot{\pi}) = 1$.
   (b) If the checks pass, it sets $b = 1$, computes locally $\texttt{Com}(\pi, r) = g_{\mathcal{C}}$, and sends $g_{\mathcal{C}}$ to the contract. Otherwise, it sets $b = 0$ and $g_{\mathcal{C}} = \perp$.
3. **Prove**. For either $\mathcal{B}$ or $\mathcal{C}$ to prove, it sends $\ddot{\pi} := (\pi, r)$ to the smart contract.
4. **Verify**. $\texttt{SAP.verify}(\ddot{\pi}, g_{\mathcal{B}}, g_{\mathcal{C}}, adr_{\mathcal{B}}, adr_{\mathcal{C}})$
   The following steps are taken by the smart contract.
   (a) Ensures $g_{\mathcal{B}}$ and $g_{\mathcal{C}}$ were sent from $adr_{\mathcal{B}}$ and $adr_{\mathcal{C}}$ respectively.
   (b) Ensures $\texttt{Ver}(g_{\mathcal{B}}, \ddot{\pi}) = \texttt{Ver}(g_{\mathcal{C}}, \ddot{\pi}) = 1$.
   (c) Outputs $s = 1$, if the checks, in steps 4a and 4b, pass. It outputs $s = 0$, otherwise.

1.6 Pseudorandom Functions

Informally, a pseudorandom function is a deterministic function that takes a key of length Λ and an input; and outputs a value indistinguishable from that of a truly random function. In this paper, we use pseudorandom function: $\texttt{PRF} : \{0,1\}^\Lambda \times \{0,1\}^* \to \mathbb{F}_p$, where p is a large prime number, $|p| = \lambda$, and (Λ, λ) are the security parameters. In practice, a pseudorandom function can be obtained from an efficient block cipher [7].

2 An Overview of Payment with Dispute Resolution Scheme

In this section, we provide an overview of a Payment with Dispute Resolution (PwDR) scheme. Simply put, a PwDR scheme allows a customer to interact with its bank (via the online banking scheme) to transfer a certain amount of money from its account to another account in a transparent manner; meanwhile, it offers a distinct feature. Namely, when an APP fraud takes place, it lets an honest customer raise a dispute and *prove* to a third-party dispute resolver that it has acted honestly, so it can be reimbursed. It offers other features too. Specifically, an honest bank can also prove it has acted honestly. PwDR lets the parties prove their innocence without their counter-party's collaboration. It also ensures the message exchanged between a bank and customer remains confidential and even the party which resolves the dispute between the two learns as little information as possible. The PwDR scheme can be considered as an extension to the existing online banking system. To design such a scheme, we need to address several challenges. The rest of this section highlights these challenges.

2.1 Challenge 1: Lack of Transparent Logs

In the current online banking system, during a payment journey, the messages exchanged between customer and bank is usually logged by the bank and is not accessible to the customer without the bank's collaboration. Even if the bank provides access to the transaction logs, there is no guarantee that the logs have remained intact. Due to the lack of a transparent logging mechanism, a customer or bank can wrongly claim that (a) it has sent a certain message or warning to its counter-party or (b) it has never received a certain message, e.g., due to hardware or software failure. Thus, it would be hard for an honest party (especially a customer) to prove its innocence. To address this challenge, the PwDR protocol uses a standard smart contract (as a public bulletin board) to which each party needs to send (a copy of) all outgoing messages, e.g., payment requests, warnings, and confirmation of payments.

2.2 Challenge 2: Lack of Effective Warning's Accurate Definition in Banking

One of the determining factors in the process of allocating liability to the customers (after an APP fraud occurs) is paying attention to and following "warning(s)", according to the "Contingent Reimbursement Model" (CRM) code [9]. However, there exists no publicly available study and evidence on the effectiveness of every warning used and provided by the bank. Therefore, we cannot hold a customer accountable for becoming the fraud's victim, even if the related warnings are ignored. To address this challenge, we let a warning's effectiveness is determined on a case-by-case basis after the fraud takes place. In particular, the protocol provides an opportunity to a victim to challenge a certain warning whose effectiveness will be assessed by a committee, i.e., a set of arbiters. In this setting, each member of the committee provides (an encoding of) its verdict to the smart contract, from which a dispute resolver retrieves all verdicts to find out the final verdict.

2.3 Challenge 3: Linking Off-chain Payments with a Smart Contract

Recall that an APP fraud occurs when a payment is made. In the case where a bank sends a message (to the smart contract) to claim that it has transferred the money following the customer's request, it is not possible to automatically validate such a claim (e.g., for the smart contract doe the check) as the money transfer takes place outside of the blockchain network. To address this challenge, the protocol lets a customer raise a dispute and report it to the smart contract when it detects an inconsistency (e.g., the bank did not transfer the money but it wrongly declared it did so, or when it transferred the money but did not declare it). In this case, the above committee members investigate and provide their verdicts to the smart contract that allows the dispute resolver to extract the final verdict.

2.4 Challenge 4: Preserving Privacy

Although the use of a public transparent logging mechanism plays a vital role in resolving disputes, if no privacy-preserving mechanism is used, then it can violate parties' privacy, e.g., the customers' payment detail, bank's messages to the customer, or even each arbiter's verdict. To protect the privacy of the bank's and customers' messages against the public (and other customers), the PwDR protocol lets the customer and bank provably agree on encoding-decoding tokens that let them to encode their outgoing messages (sent to the smart contract). Later, either party can provide the token to a third party which can independently check the tokens' correctness, and decode the messages. To protect the privacy of the committee members' verdicts from the dispute resolver, the PwDR protocol ensures that the dispute resolver can learn only the final verdict without being able to link a verdict to a specific member of the committee.

3 Definition of Payment with Dispute Resolution Scheme

In this section, we present a formal definition of a PwDR scheme. We first provide the scheme's syntax. Then, we formally define its correctness and security properties.

Definition 2. *A payment with dispute resolution scheme includes the following processes PwDR* := (`keyGen`, `bankInit`, `customerInit`, `genUpdateRequest`, `insertNewPayee`, `genWarning`, `genPaymentRequest`, `makePayment`, `genComplaint`, `verComplaint`, `resDispute`)*. It involves six types of entities; namely, bank $\mathcal{B}$, customer $\mathcal{C}$, smart contract $\mathcal{S}$, certificate generator $\mathcal{G}$, set of arbiters $\mathcal{D} : \{\mathcal{D}_1, ..., \mathcal{D}_n\}$, and a dispute resolver $\mathcal{DR}$. The processes of PwDR are defined below.*

- `keyGen`$(1^\lambda) \rightarrow (sk, pk)$. *It is a probabilistic algorithm run independently by $\mathcal{G}$ and one of the arbiters, $\mathcal{D}_j$. It takes as input a security parameter 1^λ. It outputs a pair of secret keys $sk := (sk_\mathcal{G}, sk_\mathcal{D})$ and public keys $pk := (pk_\mathcal{G}, pk_\mathcal{D})$. The public key pair, pk, is sent to all participants.*

- `bankInit`$(1^\lambda) \rightarrow (T, pp, \boldsymbol{l})$. *It is run by $\mathcal{B}$. It takes as input security parameter 1^λ. It allocates private parameters to $\ddot{\pi}_1$ and $\ddot{\pi}_2$. It generates an encoding-decoding token T, where $T := (T_1, T_2)$, each T_i contains a secret value $\ddot{\pi}_i$ and its public witness g_i. Given a value and its witness anyone can check if they match. It also generates a set of (additional) public parameters, pp, one of which is e that is a threshold parameter. It also generates an empty list, $\boldsymbol{l}$. It outputs T, pp and $\boldsymbol{l}$. $\mathcal{B}$ sends $(\ddot{\pi}_1, \ddot{\pi}_2)$ to $\mathcal{C}$ and sends $(g_1, g_2, pp, \boldsymbol{l})$ to $\mathcal{S}$.*

- `customerInit`$(1^\lambda, T, pp) \rightarrow a$. *It is a deterministic algorithm run by $\mathcal{C}$. It takes as input security parameter 1^λ, token T, and public parameters pp. It checks the correctness of the elements in T and pp. If the checks pass, it outputs 1. Otherwise, it outputs 0.*

- `genUpdateRequest`$(T, f, \boldsymbol{l}) \rightarrow \hat{m}_1^{(\mathcal{C})}$. *It is a deterministic process run by $\mathcal{C}$. It takes as input token T, new payee's detail f and empty payees' list $\boldsymbol{l}$. It generates $m_1^{(\mathcal{C})}$ which is an update request to the payees' list. It uses $T_1 \in T$ to encode $m_1^{(\mathcal{C})}$ which results in $\hat{m}_1^{(\mathcal{C})}$. It outputs $\hat{m}_1^{(\mathcal{C})}$. $\mathcal{C}$ sends the output to $\mathcal{S}$.*

- `insertNewPayee`$(\hat{m}_1^{(\mathcal{C})}, \boldsymbol{l}) \rightarrow \hat{\boldsymbol{l}}$. *It is a deterministic algorithm run by $\mathcal{S}$. It takes as input $\mathcal{C}$'s encoded update request $\hat{m}_1^{(\mathcal{C})}$, and $\mathcal{C}$'s payees' list $\boldsymbol{l}$. It inserts the new payee's detail into $\boldsymbol{l}$ and returns an updated list, $\hat{\boldsymbol{l}}$.*

- `genWarning`$(T, \hat{\boldsymbol{l}}, aux) \rightarrow \hat{m}_1^{(\mathcal{B})}$. *It is run by $\mathcal{B}$. It takes as input token T, $\mathcal{C}$'s encoded payees' list $\hat{\boldsymbol{l}}$ and auxiliary information: aux, e.g., a set of policies. Using $T_1 \in T$, it decodes and checks all elements of the list, e.g., whether they comply with the policies. If the check passes, then it sets $m_1^{(\mathcal{B})}$ = "pass"; otherwise, it sets $m_1^{(\mathcal{B})} = warning$, where the warning is a string containing a warning detail along with the string "warning". It uses T_1 to encode $m_1^{(\mathcal{B})}$ which yields $\hat{m}_1^{(\mathcal{B})}$. It outputs $\hat{m}_1^{(\mathcal{B})}$. $\mathcal{B}$ sends $\hat{m}_1^{(\mathcal{B})}$ to $\mathcal{S}$.*

- `genPaymentRequest`$(T, in_f, \hat{\boldsymbol{l}}, \hat{m}_1^{(\mathcal{B})}) \rightarrow \hat{m}_2^{(\mathcal{C})}$. *It is run by $\mathcal{C}$. It takes as input token T, a payment detail in_f, encoded payees' list $\hat{\boldsymbol{l}}$, and encoded warning message, $\hat{m}_1^{(\mathcal{B})}$. Using $T_1 \in T$, it decodes $\hat{\boldsymbol{l}}$ and $\hat{m}_1^{(\mathcal{B})}$ yielding $\boldsymbol{l}$ and $m_1^{(\mathcal{B})}$ respectively. It checks the warning. It sets $m_2^{(\mathcal{C})} = \phi$, if it does not want to proceed. Otherwise, it sets $m_2^{(\mathcal{C})}$ according to the content of in_f and $\boldsymbol{l}$ (e.g., the amount of payment and payee's detail). It uses T_1 to encode $m_2^{(\mathcal{C})}$ resulting in $\hat{m}_2^{(\mathcal{C})}$. It outputs $\hat{m}_2^{(\mathcal{C})}$. $\mathcal{C}$ sends $\hat{m}_2^{(\mathcal{C})}$ to $\mathcal{S}$.*

- `makePayment`$(T, \hat{m}_2^{(\mathcal{C})}) \rightarrow \hat{m}_2^{(\mathcal{B})}$. *It is a deterministic process run by $\mathcal{B}$. It takes as input token T, and encoded payment detail $\hat{m}_2^{(\mathcal{C})}$. Using $T_1 \in T$, it decodes $\hat{m}_2^{(\mathcal{C})}$ and checks the result's validity, e.g., ensures it is well-formed or $\mathcal{C}$ has enough credit. If the check passes, then it makes the payment and sets $m_2^{(\mathcal{B})}$ = "paid". Otherwise, it sets $m_2^{(\mathcal{B})} = \phi$. It uses T_1 to encode $m_2^{(\mathcal{B})}$ yielding $\hat{m}_2^{(\mathcal{B})}$. It outputs $\hat{m}_2^{(\mathcal{B})}$. $\mathcal{B}$ sends $\hat{m}_2^{(\mathcal{B})}$ to $\mathcal{S}$.*

- `genComplaint`$(\hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}, T, pk, aux_f) \rightarrow (\hat{z}, \hat{\ddot{\pi}})$. *It is run by $\mathcal{C}$. It takes as input the encoded warning message $\hat{m}_1^{(\mathcal{B})}$, encoded payment message $\hat{m}_2^{(\mathcal{B})}$, token T, public key pk, and auxiliary information aux_f. It initially sets fresh strings (z_1, z_2, z_3) to null. Using $T_1 \in T$, it decodes $\hat{m}_1^{(\mathcal{B})}$ and $\hat{m}_2^{(\mathcal{B})}$ and checks the results' content. If it wants to complain that (i) "pass" message should have been a warning or (ii) no message was provided, it sets z_1 to "challenge message". If its complaint is about the warning's effectiveness, it sets z_2 to a combination of an evidence $u \in aux_f$, the evidence's certificate $sig \in aux_f$, the certificate's public parameter, and "challenge warning", where the certificate is obtained from $\mathcal{G}$ via a query, Q. In certain cases, the certificate might be empty. If its complaint is about the payment, it sets z_3 to "challenge payment". It uses T_1 to encode $z := (z_1, z_2, z_3)$ and uses $pk_\mathcal{D}$ to encode $\ddot{\pi} := (\ddot{\pi}_1, \ddot{\pi}_2) \in T$. This results in $\hat{z}$ and $\hat{\ddot{\pi}}$ respectively. It outputs $(\hat{z}, \hat{\ddot{\pi}})$. $\mathcal{C}$ sends the pair to $\mathcal{S}$.*

- $\texttt{verComplaint}(\hat{z}, \hat{\ddot{\pi}}, g, \hat{\boldsymbol{m}}, \hat{\boldsymbol{l}}, j, sk_{\mathcal{D}}, aux, pp) \rightarrow \hat{\boldsymbol{w}}_j$. *It is run by every arbiter $\mathcal{D}_j$. It takes as input $\mathcal{C}$'s encoded complaint $\hat{z}$, encoded private parameters $\hat{\ddot{\pi}}$, the tokens' public parameters $g := (g_1, g_2)$, encoded messages $\hat{\boldsymbol{m}} = [\hat{m}_1^{(\mathcal{C})}, \hat{m}_2^{(\mathcal{C})}, \hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}]$, encoded payees' list $\hat{\boldsymbol{l}}$, the arbiter's index j, secret key $sk_{\mathcal{D}}$, auxiliary information aux, and public parameters pp. It uses $sk_{\mathcal{D}}$ to decode $\hat{\ddot{\pi}}$ that yields $\ddot{\pi} := (\ddot{\pi}_1, \ddot{\pi}_2)$. It uses $\ddot{\pi}_1$ to decode $\hat{z}, \hat{\boldsymbol{m}}$, and $\hat{\boldsymbol{l}}$ that results in $z := (z_1, z_2, z_3), \boldsymbol{m}$, and $\boldsymbol{l}$ respectively. It checks if $\ddot{\pi}_i$ matches g_i. If the check fails, it aborts; otherwise, it continues. It checks if $m_1^{(\mathcal{C})}$ and $m_2^{(\mathcal{C})}$ are non-empty; it aborts if the checks fail. It sets fresh parameters $(w_{1,j}, w_{2,j}, w_{3,j}, w_{4,j})$ to ϕ. If "challenge message" $\in z_1$, given $\boldsymbol{l}$, it checks whether "pass" message (in $m_1^{(\mathcal{B})}$) was given correctly or the missing message did not play any role in preventing the scam. If either checks passes, it sets $w_{1,j} = 1$; otherwise, it sets $w_{1,j} = 0$. If "challenge warning" $\in z_2$, it verifies the certificate in z_2. If it is invalid, it sets $w_{3,j} = 1$. If it is valid, it sets $w_{3,j} = 0$. It determines the effectiveness of the warning (in $m_1^{(\mathcal{B})}$), by running a process which determines that, i.e., $\texttt{checkWarning}(.) \in aux$. If it is effective, i.e., $\texttt{checkWarning}(m_1^{(\mathcal{B})}) = 1$, it sets its verdict to 1, i.e., $w_{2,j} = 1$; otherwise, it sets $w_{2,j} = 0$. If "challenge payment" $\in z_3$, it checks if the payment was made (with the help of $m_2^{(\mathcal{B})}$). If the check passes, it sets $w_{4,j} = 0$; otherwise, it sets $w_{4,j} = 1$. It uses $\ddot{\pi}_2$ to encode $\boldsymbol{w}_j = [w_{1,j}, w_{2,j}, w_{3,j}, w_{4,j}]$ yielding $\hat{\boldsymbol{w}}_j = [\hat{w}_{1,j}, \hat{w}_{2,j}, \hat{w}_{3,j}, \hat{w}_{4,j}]$. It outputs $\hat{\boldsymbol{w}}_j$. $\mathcal{D}_j$ sends $\hat{\boldsymbol{w}}_j$ to $\mathcal{S}$.*
- $\texttt{resDispute}(T_2, \hat{\boldsymbol{w}}, pp) \rightarrow \boldsymbol{v}$. *It is a deterministic algorithm run by $\mathcal{DR}$. It takes as input token T_2, arbiters' encoded verdicts $\hat{\boldsymbol{w}} = [\hat{\boldsymbol{w}}_1, ..., \hat{\boldsymbol{w}}_n]$, and public parameters pp. It checks if the token's parameters match. If the check fails, it aborts; otherwise, it proceeds. It uses $\ddot{\pi}_2 \in T_2$ to decode $\hat{\boldsymbol{w}}$ and from the result it extracts final verdicts $\boldsymbol{v} = [v_1, ..., v_4]$. The extraction procedure ensures each v_i is set to 0 only if at least e arbiters' verdicts (i.e., $w_{i,j,}$) is 0, where $e \in pp$. It outputs $\boldsymbol{v}$. Note, if $v_4 = 0$ and (i) either $v_1 = 0$ (ii) or $v_2 = 0$ and $v_3 = 0$, then customer $\mathcal{C}$ must be reimbursed.*

Informally, a PwDR scheme has two properties; namely, *correctness* and *security*. Correctness requires that (in the absence of a fraudster) the payment journey is completed without the need for (i) the honest customer to complain and (ii) the honest bank to reimburse the customer. Below, we formally state it.

Definition 3 (Correctness). *A PwDR scheme is correct if the key generation algorithm produces keys* $\texttt{keyGen}(1^\lambda) \rightarrow (sk, pk)$ *such that for any payee's detail f, payment's detail in_f, and auxiliary information (aux, aux_f), if* $\texttt{bankInit}(1^\lambda) \rightarrow (T, pp, \boldsymbol{l})$, $\texttt{customerInit}(1^\lambda, T, pp) \rightarrow a$, $\texttt{genUpdateRequest}(T, f, \boldsymbol{l}) \rightarrow \hat{m}_1^{(\mathcal{C})}$, $\texttt{insertNewPayee}(\hat{m}_1^{(\mathcal{C})}, \boldsymbol{l}) \rightarrow \hat{\boldsymbol{l}}$, $\texttt{genWarning}(T, \hat{\boldsymbol{l}}, aux) \rightarrow \hat{m}_1^{(\mathcal{B})}$, $\texttt{genPaymentRequest}(T, in_f, \hat{\boldsymbol{l}}, \hat{m}_1^{(\mathcal{B})}) \rightarrow \hat{m}_2^{(\mathcal{C})}$, $\texttt{makePayment}(T, \hat{m}_2^{(\mathcal{C})}) \rightarrow \hat{m}_2^{(\mathcal{B})}$, $\texttt{genComplaint}(\hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}, T, pk, aux_f) \rightarrow (\hat{z}, \hat{\ddot{\pi}})$, $\forall j \in [n]$: $\Big(\texttt{verComplaint}(\hat{z}, \hat{\ddot{\pi}}, g, \hat{\boldsymbol{m}}, \hat{\boldsymbol{l}}, j, sk_{\mathcal{D}}, aux, pp) \rightarrow \hat{\boldsymbol{w}}_j\Big)$, $\texttt{resDispute}(T_2, \hat{\boldsymbol{w}}, pp) \rightarrow \boldsymbol{v}$, *then* $(z_1 = z_2 = z_3 = \phi) \wedge (\boldsymbol{v} = \phi)$, *where* $g := (g_1, g_2) \in T$, $\hat{\boldsymbol{m}} = [\hat{m}_1^{(\mathcal{C})}, \hat{m}_2^{(\mathcal{C})}, \hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}]$, $\hat{\boldsymbol{w}} = [\hat{\boldsymbol{w}}_1, ..., \hat{\boldsymbol{w}}_n]$, *and* $z := (z_1, z_2, z_3)$ *is the decoded* $\hat{z}$.

A PwDR scheme is secure if it satisfies three main properties; namely, (a) security against a malicious victim, (b) security against a malicious bank, and (c) privacy. Below, we formally define them. Intuitively, security against a malicious victim requires that the victim of the APP scam which is not qualified for the reimbursement should not be reimbursed (despite it tries to be). More specifically, a corrupt victim cannot (a) make committee members, $\mathcal{D}_j$s, conclude that $\mathcal{B}$ should have provided a warning, although $\mathcal{B}$ has done so, or (b) make $\mathcal{DR}$ conclude that the "pass" message was incorrectly given or a vital warning message was missing despite all $\mathcal{D}_j$s do not believe so, or (c) persuade $\mathcal{D}_j$s to conclude that the warning was ineffective although it was effective, or (d) make $\mathcal{DR}$ believe that the warning message was ineffective although all $\mathcal{D}_j$s do not believe that, or (e) convince $\mathcal{D}_j$s to accept an invalid certificate, or (f) make $\mathcal{DR}$ believe that all $\mathcal{D}_j$s accepted the certificate although they did not, except for a negligible probability.

Definition 4 (Security against a malicious victim). *A PwDR scheme is secure against a malicious victim, if for any security parameter λ, auxiliary information aux, and probabilistic polynomial-time adversary $\mathcal{A}$, there exists a negligible function $\mu(\cdot)$, such that for an experiment $\mathsf{Exp}_1^{\mathcal{A}}$:*

$\mathsf{Exp}_1^{\mathcal{A}}(1^\lambda,\ aux)$

$\texttt{keyGen}(1^\lambda) \rightarrow (sk, pk)$
$\texttt{bankInit}(1^\lambda) \rightarrow (T, pp, \boldsymbol{l})$
$\mathcal{A}(1^\lambda, T, pp, \boldsymbol{l}) \rightarrow \hat{m}_1^{(\mathcal{C})}$
$\texttt{insertNewPayee}(\hat{m}_1^{(\mathcal{C})}, \boldsymbol{l}) \rightarrow \hat{\boldsymbol{l}}$
$\texttt{genWarning}(T, \hat{\boldsymbol{l}}, aux) \rightarrow \hat{m}_1^{(\mathcal{B})}$
$\mathcal{A}(T, \hat{\boldsymbol{l}}, \hat{m}_1^{(\mathcal{B})}) \rightarrow \hat{m}_2^{(\mathcal{C})}$
$\texttt{makePayment}(T, \hat{m}_2^{(\mathcal{C})}) \rightarrow \hat{m}_2^{(\mathcal{B})}$
$\mathcal{A}(\hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}, T, pk) \rightarrow (\hat{z}, \hat{\hat{\pi}})$
$\forall j, j \in [n]:$
$\Big(\texttt{verComplaint}(\hat{z}, \hat{\hat{\pi}}, g, \hat{\boldsymbol{m}}, \hat{\boldsymbol{l}}, j, sk_{\mathcal{D}}, aux, pp) \rightarrow \hat{\boldsymbol{w}}_j = [\hat{w}_{1,j}, \hat{w}_{2,j}, \hat{w}_{3,j}, \hat{w}_{4,j}]\Big)$
$\texttt{resDispute}(T_2, \hat{\boldsymbol{w}}, pp) \rightarrow \boldsymbol{v} = [v_1, ..., v_4]$

it holds that:

$$\Pr\left[\begin{array}{l}\Big((m_1^{(\mathcal{B})} = warning) \wedge (\sum_{j=1}^{n} w_{1,j} < n)\Big) \\ \vee\Big((\sum_{j=1}^{n} w_{1,j} = n) \wedge (v_1 = 0)\Big) \\ \vee\Big((\texttt{checkWarning}(m_1^{(\mathcal{B})}) = 1) \wedge (\sum_{j=1}^{n} w_{2,j} < n)\Big) \\ \vee\Big((\sum_{j=1}^{n} w_{2,j} = n) \wedge (v_2 = 0)\Big) \\ \vee\Big(u \notin Q \wedge \texttt{Sig.ver}(pk, u, sig) = 1\Big) \\ \vee\Big((\sum_{j=1}^{n} w_{3,j} < n) \wedge (v_3 = 0)\Big)\end{array} : \mathsf{Exp}_1^{\mathcal{A}}(\mathsf{input})\right] \leq \mu(\lambda),$$

where $g := (g_1, g_2) \in T$, $\hat{\boldsymbol{m}} = [\hat{m}_1^{(\mathcal{C})}, \hat{m}_2^{(\mathcal{C})}, \hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}]$, $(w_{1,j}, .., w_{3,j})$ *are the result of decoding* $(\hat{w}_{1,j}, .., \hat{w}_{3,j}) \in \hat{\boldsymbol{w}}$, $\texttt{checkWarning}(.)$ *determines a warning's effectiveness,* $\mathsf{input} := (1^\lambda, aux)$, $(u, sig) \in x$, $sk_{\mathcal{D}} \in sk$, *and* n *is the total number of arbiters. The probability is taken over the uniform choice of* sk, *randomness used in the blockchain's primitives (e.g., in signatures), randomness used during the encoding, and the randomness of* $\mathcal{A}$.

Intuitively, security against a malicious bank requires that a malicious bank should not be able to disqualify an honest victim of the APP scam from being reimbursed. In particular, a corrupt bank cannot (a) make $\mathcal{DR}$ conclude that the "pass" message was correctly given or an important warning message was not missing despite not all $\mathcal{D}_j$s believe so, or (b) convince $\mathcal{DR}$ that the warning message was effective although not all $\mathcal{D}_j$s believe so, or (c) make $\mathcal{DR}$ believe that all $\mathcal{D}_j$s did not accept the certificate although they did, or (d) make $\mathcal{DR}$ believe that no payment was made, although all $\mathcal{D}_j$s believe the opposite, except for a negligible probability.

Definition 5 (Security against a malicious bank). *A PwDR scheme is secure against a malicious bank, if for any security parameter* λ, *auxiliary information* aux, *and probabilistic polynomial-time adversary* $\mathcal{A}$, *there exists a negligible function* $\mu(\cdot)$, *such that for an experiment* $\mathsf{Exp}_2^{\mathcal{A}}$*:*

$\mathsf{Exp}_2^{\mathcal{A}}(1^\lambda,\ aux)$

$\texttt{keyGen}(1^\lambda) \rightarrow (sk, pk)$
$\mathcal{A}(1^\lambda) \rightarrow (T, pp, \boldsymbol{l}, f, in_f, aux_f)$
$\texttt{customerInit}(1^\lambda, T, pp) \rightarrow a$
$\texttt{genUpdateRequest}(T, f, \boldsymbol{l}) \rightarrow \hat{m}_1^{(\mathcal{C})}$
$\texttt{insertNewPayee}(\hat{m}_1^{(\mathcal{C})}, \boldsymbol{l}) \rightarrow \hat{\boldsymbol{l}}$
$\mathcal{A}(T, \hat{\boldsymbol{l}}, aux) \rightarrow \hat{m}_1^{(\mathcal{B})}$
$\texttt{genPaymentRequest}(T, in_f, \hat{\boldsymbol{l}}, \hat{m}_1^{(\mathcal{B})}) \rightarrow \hat{m}_2^{(\mathcal{C})}$
$\mathcal{A}(T, \hat{m}_2^{(\mathcal{C})}) \rightarrow \hat{m}_2^{(\mathcal{B})}$
$\texttt{genComplaint}(\hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}, T, pk, aux_f) \rightarrow (\hat{z}, \hat{\hat{\pi}})$
$\forall j, j \in [n]:$
$\Big(\texttt{verComplaint}(\hat{z}, \hat{\hat{\pi}}, g, \hat{\boldsymbol{m}}, \hat{\boldsymbol{l}}, j, sk_{\mathcal{D}}, aux, pp) \rightarrow \hat{\boldsymbol{w}}_j = [\hat{w}_{1,j}, \hat{w}_{2,j}, \hat{w}_{3,j}, \hat{w}_{4,j}]\Big)$
$\texttt{resDispute}(T_2, \hat{\boldsymbol{w}}, pp) \rightarrow \boldsymbol{v} = [v_1, ..., v_4]$

it holds that:

$$\Pr\left[\begin{array}{l}\Big((\sum_{j=1}^{n} w_{1,j} < n) \wedge (v_1 = 1)\Big) \\ \vee\Big((\sum_{j=1}^{n} w_{2,j} < n) \wedge (v_2 = 1)\Big) \\ \vee\Big((\sum_{j=1}^{n} w_{3,j} = 0) \wedge (v_3 = 1)\Big) \\ \vee\Big((\sum_{j=1}^{n} w_{4,j} = 0) \wedge (v_4 = 1)\Big)\end{array} : \mathsf{Exp}_2^{\mathcal{A}}(\mathsf{input})\right] \leq \mu(\lambda),$$

where $g := (g_1, g_2) \in T$, $\hat{\boldsymbol{m}} = [\hat{m}_1^{(\mathcal{C})}, \hat{m}_2^{(\mathcal{C})}, \hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}]$, $(w_{1,j}, .., w_{3,j})$ are the result of decoding $(\hat{w}_{1,j}, .., \hat{w}_{3,j}) \in \hat{\boldsymbol{w}}$, $\mathsf{input} := (1^\lambda, aux)$, $sk_{\mathcal{D}} \in sk$, n is the total number of arbiters. The probability is taken over the uniform choice of sk, randomness used in the blockchain's primitives, randomness used during the encoding, and the randomness of $\mathcal{A}$.

Now we move on to privacy. Informally, a PwDR is privacy-preserving if it protects the privacy of (1) the customers', bank's, and arbiters' messages (except public parameters) from non-participants of the protocol, including other customers, and (2) each individual arbiter's verdict from a third-party dispute resolver which sees the final verdict.

Definition 6 (Privacy). *A PwDR preserves privacy if the following two properties are satisfied.*

1. *For any probabilistic polynomial-time adversary $\mathcal{A}_1$, security parameter λ, and auxiliary information aux, there exists a negligible function $\mu(\cdot)$, such that for any experiment $\mathsf{Exp}_3^{\mathcal{A}_1}$:*

$\mathsf{Exp}_3^{\mathcal{A}_1}(1^\lambda,\ aux)$

$\texttt{keyGen}(1^\lambda) \rightarrow (sk, pk)$
$\texttt{bankInit}(1^\lambda) \rightarrow (T, pp, \boldsymbol{l})$
$\texttt{customerInit}(1^\lambda, T, pp) \rightarrow a$
$\mathcal{A}_1(1^\lambda, pk, a, pp, g, \boldsymbol{l}) \rightarrow \Big((f_0, f_1), (in_{f_0}, in_{f_1}), (aux_{f_0}, aux_{f_1})\Big)$
$\gamma \xleftarrow{\$} \{0,1\}$
$\texttt{genUpdateRequest}(T, f_\gamma, \boldsymbol{l}) \rightarrow \hat{m}_1^{(\mathcal{C})}$
$\texttt{insertNewPayee}(\hat{m}_1^{(\mathcal{C})}, \boldsymbol{l}) \rightarrow \hat{\boldsymbol{l}}$
$\texttt{genWarning}(T, \hat{\boldsymbol{l}}, aux) \rightarrow \hat{m}_1^{(\mathcal{B})}$
$\texttt{genPaymentRequest}(T, in_{f_\gamma}, \hat{\boldsymbol{l}}, \hat{m}_1^{(\mathcal{B})}) \rightarrow \hat{m}_2^{(\mathcal{C})}$
$\texttt{makePayment}(T, \hat{m}_2^{(\mathcal{C})}) \rightarrow \hat{m}_2^{(\mathcal{B})}$
$\texttt{genComplaint}(\hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}, T, pk, aux_{f_\gamma}) \rightarrow (\hat{z}, \hat{\hat{\pi}})$
$\forall j, j \in [n]:$
$\Big(\texttt{verComplaint}(\hat{z}, \hat{\hat{\pi}}, g, \hat{\boldsymbol{m}}, \hat{\boldsymbol{l}}, j, sk_{\mathcal{D}}, aux, pp) \rightarrow \hat{\boldsymbol{w}}_j\Big)$
$\texttt{resDispute}(T_2, \hat{\boldsymbol{w}}, pp) \rightarrow \boldsymbol{v}$

it holds that:

$$\Pr\left[\mathcal{A}_1(\hat{\boldsymbol{m}}, \hat{\boldsymbol{l}}, \hat{z}, \hat{\hat{\pi}}, \hat{\boldsymbol{w}}) \rightarrow \gamma \,:\, \mathsf{Exp}_3^{\mathcal{A}_1}(\mathsf{input})\right] \leq \frac{1}{2} + \mu(\lambda).$$

2. *For any probabilistic polynomial-time adversaries $\mathcal{A}_2$ and $\mathcal{A}_3$, security parameter λ, and auxiliary information aux, there exists a negligible function $\mu(\cdot)$, such that for any experiment $\mathsf{Exp}_4^{\mathcal{A}_2}$:*

$\mathsf{Exp}_4^{\mathcal{A}_2}(1^\lambda,\ aux)$

$\texttt{keyGen}(1^\lambda) \rightarrow (sk, pk)$
$\texttt{bankInit}(1^\lambda) \rightarrow (T, pp, \boldsymbol{l})$
$\texttt{customerInit}(1^\lambda, T, pp) \rightarrow a$
$\mathcal{A}_2(1^\lambda, pk, a, pp, \boldsymbol{l}) \rightarrow (f, in_f, aux_f)$
$\texttt{genUpdateRequest}(T, f, \boldsymbol{l}) \rightarrow \hat{m}_1^{(\mathcal{C})}$
$\texttt{insertNewPayee}(\hat{m}_1^{(\mathcal{C})}, \boldsymbol{l}) \rightarrow \hat{\boldsymbol{l}}$
$\mathcal{A}_2(T, \hat{\boldsymbol{l}}, aux) \rightarrow \hat{m}_1^{(\mathcal{B})}$
$\texttt{genPaymentRequest}(T, in_f, \hat{\boldsymbol{l}}, \hat{m}_1^{(\mathcal{B})}) \rightarrow \hat{m}_2^{(\mathcal{C})}$
$\mathcal{A}_2(T, pk, aux_f, \hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{C})}) \rightarrow (\hat{m}_2^{(\mathcal{B})}, \hat{z}, \hat{\hat{\pi}})$
$\forall j, j \in [n]:$
$\Big(\texttt{verComplaint}(\hat{z}, \hat{\hat{\pi}}, g, \hat{\boldsymbol{m}}, \hat{\boldsymbol{l}}, j, sk_{\mathcal{D}}, aux, pp) \rightarrow \hat{\boldsymbol{w}}_j\Big)$
$\texttt{resDispute}(T_2, \hat{\boldsymbol{w}}, pp) \rightarrow \boldsymbol{v}$

it holds that:

$$\Pr\left[\mathcal{A}_3(T_2, pk, pp, g, \hat{\boldsymbol{m}}, \hat{\boldsymbol{l}}, \hat{z}, \hat{\hat{\pi}}, \hat{\boldsymbol{w}}, \boldsymbol{v}) \rightarrow w_j \,:\, \mathsf{Exp}_4^{\mathcal{A}_2}(\mathsf{input})\right] \leq Pr' + \mu(\lambda),$$

where $g := (g_1, g_2) \in T$, $\hat{\boldsymbol{m}} = [\hat{m}_1^{(\mathcal{C})}, \hat{m}_2^{(\mathcal{C})}, \hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}]$, $\hat{\boldsymbol{w}} = [\hat{\boldsymbol{w}}_1, .., \hat{\boldsymbol{w}}_n]$, $\mathsf{input} := (1^\lambda, aux)$, $sk_{\mathcal{D}} \in sk$, n is the total number of arbiters. Moreover, Pr' is defined as follows. Let arbiter $\mathcal{D}_i$ output 0 and 1 with probabilities $Pr_{i,0}$ and $Pr_{i,1}$ respectively. Then, Pr' is defined as $Max\{Pr_{1,0}, Pr_{1,1}, ..., Pr_{n,0},\ Pr_{n,1}\}$. In the above privacy definition, the probability is taken over the uniform choice of sk, the probability that

each $\mathcal{D}_j$ outputs 0 or 1, the randomness used in the blockchain's primitives, the randomness used during the encoding, and the randomness of $\mathcal{A}_1$ and $\mathcal{A}_2$.

Definition 7 (Security). *A PwDR is secure if it meets security against a malicious victim, security against a malicious bank, and preserves privacy with respect to definitions 4, 5, and 6 respectively.*

4 Payment with Dispute Resolution Protocol

In this section, first we provide an outline of the PwDR protocol. Then, we present a few subroutines that will be used in the PwDR scheme. After that, we describe the PwDR protocol in detail.

4.1 An Overview of the PwDR Protocol

In this section, we outline how the PwDR protocol works. For the sake of simplicity, we assume $\mathcal{C}$ wants to transfer a certain amount of money to a new payee. Initially, only for once, customer $\mathcal{C}$ and bank $\mathcal{B}$ agree on a smart contract $\mathcal{S}$. They also use SAP to provably agree on two private statements that include two secret keys. The keys will be used to encrypt messages sent to $\mathcal{S}$ and will be used for decryption by two separate types of party (i.e., $\mathcal{D}_j$s and $\mathcal{DR}$). When $\mathcal{C}$ wants to transfer money to a new payee, it signs into the standard online banking system. Then, it generates an update request that specifies the new payee's detail, encrypts the request, and sends the result to $\mathcal{S}$. After that, $\mathcal{B}$ decrypts and checks the request's validity, e.g., whether it meets its internal policy or the requirements of the "Confirmation of Payee" [4]. Depending on the request's content, $\mathcal{B}$ generates a pass or warning message. It encrypts the message and sends the result to $\mathcal{S}$. Then, $\mathcal{C}$ checks $\mathcal{B}$'s message and depending on the content of this message, it decides whether it wants to proceed to made payment. If it decides to do so, then it sends an encrypted payment detail to $\mathcal{S}$. This allows $\mathcal{B}$ to decrypt the message and locally transfer the amount of money specified in $\mathcal{C}$'s message. Once the money is transferred, $\mathcal{B}$ sends an encrypted "paid" message to $\mathcal{S}$.

Once $\mathcal{C}$ realises that it has fallen victim to an APP fraud, it raises a dispute. In particular, it generates an encrypted complaint that could challenge the effectiveness of the warning and/or any payment inconsistency (as explained in Section 2.3). It can also include in the complaint an evidence/certificate, e.g., to claim that it falls into the vulnerable customer category as defined in [9]. $\mathcal{C}$ encrypts the complaint. It also encrypts the secret key (under arbiters' public key) that it uses to encrypt the messages. It sends to $\mathcal{S}$ the ciphertexts along with a proof asserting the secret key's correctness. Upon receiving $\mathcal{C}$'s complaint, each committee member verifies the proof. If the verification passes, it decrypts and processes $\mathcal{C}$'s complaint to generate a (set of) verdict. Then, each committee member encodes its verdict and sends the encryption of the encoded verdict to $\mathcal{S}$. To resolve a dispute between $\mathcal{C}$ and $\mathcal{B}$, either of them can invoke $\mathcal{DR}$. To do so, they directly send to it one of the above secret keys and a proof asserting that key was generated correctly. $\mathcal{DR}$ verifies the proof. If the verification passes it locally decrypts the encrypted encoded verdicts (after retrieving them from $\mathcal{S}$) and then combines the result to find out the final verdict. If the final verdict indicates the legitimacy of $\mathcal{C}$'s complaint, then $\mathcal{C}$ must be reimbursed. Note, the verdicts are encoded in such a way that even after decrypting them, the dispute resolver cannot link a verdict to a committee member or even figure out how many 1 or 0 verdicts have been provided (except when all verdicts are 0). Shortly, we present novel verdict encoding-decoding protocols that offer the above features.

4.2 The PwDR Protocol's Subroutines

In this section, we present a few subroutines that will be called by the PwDR protocol.

Determining Bank's Message Status. As we stated earlier, in the payment journey the customer may receive a "pass" message or even nothing at all, e.g., due to a system failure. In such cases, a victim of an APP fraud may complain that if the pass or missing message was a warning message, then it would have prevented the victim from falling to the APP fraud. To assist the committee members to deal with

such complaints deterministically, we propose a process, called `verStat`. It is run locally by each committee member. It outputs 1 if a pass message was given correctly or the missing message could not prevent the scam, and outputs 0 otherwise. The process is presented in figure 1.

$\texttt{verStat}(add_{\mathcal{S}}, m^{(\mathcal{B})}, \boldsymbol{l}, \Delta, aux) \rightarrow w_1$

- *Input.* $add_{\mathcal{S}}$: the address of smart contract $\mathcal{S}$, $m^{(\mathcal{B})}$: $\mathcal{B}$'s warning message, $\boldsymbol{l}$: customer's payees' list, Δ: a time parameter, and aux: auxiliary information, e.g., bank's policy.
- *Output.* $w_1 = 1$: if the "pass" message had been given correctly or the missing message did not play any role in preventing the scam; $w_1 = 0$: otherwise.

1. reads the content of $\mathcal{S}$. It checks if $m^{(\mathcal{B})} =$"pass" or the encrypted warning message was not sent on time (i.e., never sent or sent after $t_0 + \Delta$). If one of the checks passes, it proceeds to the next step. Otherwise, it aborts.
2. checks the validity of customer's most recent payees' list $\boldsymbol{l}$, with the help of the auxiliary information, aux.
   - if $\boldsymbol{l}$ contains an invalid element, it sets $w_1 = 0$.
   - otherwise, it sets $w_1 = 1$.
3. returns w_1.

Fig. 1: Process to Determine a Bank's Message Status

Checking Warning's Effectiveness. To help the committee members deterministically and accurately compile a victim's complaint about the effectiveness of a warning (that bank provides during the payment journey), we propose a process, called `checkWarning`. This process is run locally by each committee member. It also allows the victims to provide a certificate/evidence as part of their complaints. The process outputs a pair (w_2, w_3). It sets $w_2 = 1$ if the given warning message is effective, and sets $w_2 = 0$, if it is not. It sets $w_3 = 1$ if the certificate that the victim provided is valid (or empty) and sets $w_3 = 0$ if it is invalid. This process is presented in figure 2.

Verdict Encoding-Decoding Protocols. Here, we present two efficient verdict encoding and decoding protocols; namely, Private Verdict Encoding (PVE) and Final Verdict Decoding (FVD) protocols. Their goal is to let a third party $\mathcal{I}$, e.g., $\mathcal{DR}$, find out whether at least one arbiter voted 0, while satisfying the following requirements. The protocols should (1) generate unlinkable verdicts, (2) not require arbiters to interact with each other for each customer, and (3) be efficient. Since, the second and third requirements are self-explanatory, we only explain the first one. Informally, the first property requires that the protocols should generate encoded verdicts and final verdict in a way that $\mathcal{I}$, given the encoded verdicts and final verdict, should not be able to (a) link a verdict to an arbiter (except when all arbiters' verdicts are 1), and (b) find out the total number of 1 or 0 verdicts when they provide different verdicts.

At a high level, the protocols work as follows. The arbiters only once for all customers agree on a secret key of a pseudorandom function. This key will allow each of them to generate a pseudorandom masking values such that if all masking values are "XOR"ed, they would cancel out each other and result 0.[1]

Each arbiter represents its verdict by (i) representing it as a parameter which is set to either 1 if the verdict is 1 or to a random value if the verdict is 0, and then (ii) masking this parameter by the above pseudorandom value. It sends the result to $\mathcal{I}$. To decode the final verdict and find out whether any arbiter

[1] This is similar to the idea used in the XOR-based secret sharing [12].

$\texttt{checkWarning}(add_{\mathcal{S}}, z, m^{(\mathcal{B})}, aux') \rightarrow (w_2, w_3)$

- *Input.* $add_{\mathcal{S}}$: the address of smart contract $\mathcal{S}$, z: $\mathcal{C}$'s complaint, $m^{(\mathcal{B})}$: $\mathcal{B}$'s warning message, and aux': auxiliary information, e.g., guideline on warnings' effectiveness.
- *Output.* $w_3 = 1$: if the certificate in z is valid or no certificate is provided; $w_3 = 0$: if the certificate is invalid. Also, $w_2 = 1$: if the given warning message is effective; $w_2 = 0$: if the warning message is ineffective.

1. parse $z = m||sig||pk||$"challenge warning". If the certificate sig is empty, then it sets $w_3 = 0$ and proceeds to step 2. Otherwise, it:
   (a) verifies the certificate: $\texttt{Sig.ver}(pk, m, sig) \rightarrow h$.
   (b) if the certificate is rejected (i.e., $h = 0$), it sets $w_3 = 1$. It goes to step 4.
   (c) otherwise (i.e., $h = 1$), it sets $w_3 = 0$ and moves onto the next step.
2. checks if "warning" $\in m^{(\mathcal{B})}$. If the check is passed, it proceeds to the next step. Otherwise, it aborts.
3. checks the warning's effectiveness, with the assistance of the evidence m and auxiliary information aux'.
   - if it is effective, it sets $w_2 = 1$.
   - otherwise, it sets $w_2 = 0$.
4. returns (w_2, w_3).

Fig. 2: Process to Check Warning's Effectiveness

voted 0, $\mathcal{I}$ does XOR all encoded verdicts. This removes the masks and XORs are verdicts' representations. If the result is 0, then all arbiters must have voted 1; therefore, the final verdict is 1. However, if the result is not 0 (i.e., a random value), then at least one of the arbiters voted 0, so the final verdict is 0. We present the encoding and decoding protocols in figures 3 and 4 respectively.

Not that the protocols' correctness holds, except a negligible probability. In particular, if two arbiters represent their verdict by an identical random value, then when they are XORed they would cannel out each other which can affect the result's correctness. The same holds if the XOR of multiple verdicts' representations results in a value that can cancel out another verdict's representation. Nevertheless, the probability that such an event occurs is negligible in the security parameter, i.e., at most $\frac{1}{2^{|p|}}$. It is evident that PVE and FVD protocols meet properties (2) and (3). The primary reason they also meet property (1) is that each masked verdict reveals nothing about the verdict (and its representation) and given the final verdict, $\mathcal{I}$ cannot distinguish between the case where there is exactly one arbiter that voted 0 and the case where multiple arbiters voted 0, as in both cases $\mathcal{I}$ extracts only a single random value, which reveals nothing about the number of arbiters which voted 0 or 1.

4.3 The PwDR Protocol

In this section, we present the PwDR protocol in detail.

1. *Generating Certificate Parameters*. $\texttt{keyGen}(1^\lambda) \rightarrow (sk, pk)$
   The certificate generator takes step 1a and an arbiter takes step 1b below.
   (a) calls $\texttt{Sig.keyGen}(1^\lambda) \rightarrow (sk_{\mathcal{G}}, pk_{\mathcal{G}})$ to generate signing secret key $sk_{\mathcal{G}}$ and verfiying public key $pk_{\mathcal{G}}$. It publishes the public key, $pk_{\mathcal{G}}$.
   (b) calls $\tilde{\texttt{keyGen}}(1^\lambda) \rightarrow (sk_{\mathcal{D}}, pk_{\mathcal{D}})$ to generate decrypting secret key $sk_{\mathcal{D}}$ and encrypting public key $pk_{\mathcal{D}}$. It publishes the public key $pk_{\mathcal{D}}$ and sends $sk_{\mathcal{D}}$ to the rest of arbiters.
   Let $sk := (sk_{\mathcal{G}}, sk_{\mathcal{D}})$ and $pk := (pk_{\mathcal{G}}, pk_{\mathcal{D}})$. Note, this phase takes place only once for all customers.
2. *Bank-side Initiation*. $\texttt{bankInit}(1^\lambda) \rightarrow (T, pp, \boldsymbol{l})$
   $\mathcal{B}$ takes the following steps.

$\texttt{PVE}(\bar{k}_0, \text{ID}, w_j, o, n, j) \rightarrow \bar{w}_j$

- *Input.* $\bar{k}_0$: a key of pseudorandom function $\texttt{PRF}(.)$, ID: a unique identifier, w_j: a verdict, o: an offset, n: the total number of arbiters, and j: an arbiter's index.
- *Output.* $\bar{w}_j$: an encoded verdict.

Arbiter $\mathcal{D}_j$ takes the following steps.

1. computes a pseudorandom value, as follows.
   - if $j < n : r_j = \texttt{PRF}(\bar{k}_0, o||j||\text{ID})$.
   - if $j = n : r_j = \bigoplus_{j=1}^{n-1} r_j$.
2. sets a fresh parameter, w'_j, as below.

$$w'_j = \begin{cases} 1, & \text{if } w_j = 1 \\ \alpha_j \xleftarrow{\$} \mathbb{F}_p, & \text{if } w_j = 0 \end{cases}$$

3. encodes w'_j as follows. $\bar{w}_j = w'_j \oplus r_j$.
4. outputs $\bar{w}_j$.

Fig. 3: Private Verdict Encoding (PVE) Protocol

$\texttt{FVD}(n, \bar{\boldsymbol{w}}) \rightarrow v$

- *Input.* n: the total number of arbiters, and $\bar{\boldsymbol{w}} = [\bar{w}_1, ..., \bar{w}_n]$: a vector of all arbiters' encodes verdicts.
- *Output.* v: final verdict.

A third-party $\mathcal{I}$ takes the following steps.

1. combines all arbiters' encoded verdicts, $\bar{w}_j \in \bar{\boldsymbol{w}}$, as follows. $c = \bigoplus_{j=1}^{n} \bar{w}_j$
2. if n is odd, then sets $c = c \oplus 1$.
3. sets the final verdict v depending on the content of c. Specifically,

$$v = \begin{cases} 1, & \text{if } c = 0 \\ 0, & \text{otherwise} \end{cases}$$

4. outputs v.

Fig. 4: Final Verdict Decoding (FVD) Protocol

(a) picks secret keys $\bar{k}_1$ and $\bar{k}_2$ for symmetric key encryption scheme and pseudorandom function $\texttt{PRF}$ respectively. It sets two private statements as $\pi_1 = \bar{k}_1$ and $\pi_2 = \bar{k}_2$.

(b) calls $\texttt{SAP.init}(1^\lambda, adr_{\mathcal{B}}, adr_{\mathcal{C}}, \pi_i) \rightarrow (r_i, g_i, adr_{\text{SAP}})$ to initiate agreements on statements $\pi_i \in \{\pi_1, \pi_2\}$ with customer $\mathcal{C}$. Let $T_i := (\ddot{\pi}_i, g_i)$ and $T := (T_1, T_2)$, where $\ddot{\pi}_i := (\pi_i, r_i)$ is the opening of g_i. It also sets parameter Δ as a time window between two specific time points, i.e., $\Delta = t_i - t_{i-1}$. Briefly, it is used to impose an upper bound on a message delay.

(c) sends $\ddot{\pi} := (\ddot{\pi}_1, \ddot{\pi}_2)$ to $\mathcal{C}$ and sends public parameter $pp := (adr_{\text{SAP}}, \Delta)$ to smart contract $\mathcal{S}$.

3. *Customer-side Initiation*. $\texttt{customerInit}(1^\lambda, T, pp) \rightarrow a$

$\mathcal{C}$ takes the following steps.

(a) calls $\texttt{SAP.agree}(\pi_i, r_i, g_i, adr_{\mathcal{B}}, adr_{\text{SAP}}) \rightarrow (g'_i, b_i)$, to check the correctness of parameters in $T_i \in T$ and (if accepted) to agree on these parameters, where $(\pi_i, r_i) \in \ddot{\pi}_i \in T_i$ and $1 \leq i \leq 2$. Note, if both $\mathcal{B}$ and $\mathcal{C}$ are honest, then $g_i = g'_i$. It also checks Δ in $\mathcal{S}$, e.g., to see whether it is sufficiently large.

(b) if the above checks fail, it sets $a = 0$. Otherwise, it sets $a = 1$. It sends a to $\mathcal{S}$.

4. *Generating Update Request*. $\texttt{genUpdateRequest}(T, f, \boldsymbol{l}) \rightarrow \hat{m}_1^{(\mathcal{C})}$
$\mathcal{C}$ takes the following steps.
   (a) sets the inputs of algorithm $\mathcal{I} \in \{\texttt{setupNewPayee}(.), \texttt{ammendExistingPayee}(.)\}$ as below.
      - if $\texttt{setupNewPayee}(.)$ is called, then it sets $m_1^{(\mathcal{C})} := (\phi, f)$, where f is new payee's detail.
      - if $\texttt{ammendExistingPayee}(.)$ is called, then it sets $m_1^{(\mathcal{C})} := (i, f)$, where i is the index of the element in $\boldsymbol{l}$ that should be changed to f.
   (b) at time t_0, sends to $\mathcal{S}$ the encryption of $m_1^{(\mathcal{C})}$, i.e., $\hat{m}_1^{(\mathcal{C})} = \texttt{Enc}(\bar{k}_1, m_1^{(\mathcal{C})})$.

5. *Inserting New Payee*. $\texttt{insertNewPayee}(\hat{m}_1^{(\mathcal{C})}, \boldsymbol{l}) \rightarrow \hat{\boldsymbol{l}}$
$\mathcal{S}$ takes the following steps.
   - if $\hat{m}_1^{(\mathcal{C})}$ is not empty, it appends $\hat{m}_1^{(\mathcal{C})}$ to the payee list $\hat{\boldsymbol{l}}$, resulting an updated list, $\hat{\boldsymbol{l}}$.
   - if $\hat{m}_1^{(\mathcal{C})}$ is empty, it does nothing.

6. *Generating Warning*. $\texttt{genWarning}(T, \hat{\boldsymbol{l}}, aux) \rightarrow \hat{m}_1^{(\mathcal{B})}$
$\mathcal{B}$ takes the following steps.
   (a) checks if the most recent list $\hat{\boldsymbol{l}}$ is not empty. If it is empty, it halts. Otherwise, it proceeds to the next step.
   (b) decrypts each element of $\hat{\boldsymbol{l}}$ and checks their correctness, e.g., checks whether each element meets its internal policy or CoP requirements stated in aux. If the check passes, it sets $m_1^{(\mathcal{B})} =$ "pass". Otherwise, it sets $m_1^{(\mathcal{B})} =$ warning, where warning is a string that contains a warning's detail concatenated with the string "warning".
   (c) at time t_1, sends to $\mathcal{S}$ the encryption of $m_1^{(\mathcal{B})}$, i.e., $\hat{m}_1^{(\mathcal{B})} = \texttt{Enc}(\bar{k}_1, m_1^{(\mathcal{B})})$.

7. *Generating Payment Request*. $\texttt{genPaymentRequest}(T, in_f, \hat{\boldsymbol{l}}, \hat{m}_1^{(\mathcal{B})}) \rightarrow \hat{m}_2^{(\mathcal{C})}$
$\mathcal{C}$ takes the following steps.
   (a) at time t_2, decrypts the content of $\hat{\boldsymbol{l}}$ and $\hat{m}_1^{(\mathcal{B})}$. It sets a payment request $m_2^{(\mathcal{C})}$ to ϕ or in_f, where in_f contains the payment's detail, e.g., the payee's detail in $\boldsymbol{l}$ and the amount it wants to transfer.
   (b) at time t_3, sends to $\mathcal{S}$ the encryption of $m_2^{(\mathcal{C})}$, i.e., $\hat{m}_2^{(\mathcal{C})} = \texttt{Enc}(\bar{k}_1, m_2^{(\mathcal{C})})$.

8. *Making Payment*. $\texttt{makePayment}(T, \hat{m}_2^{(\mathcal{C})}) \rightarrow \hat{m}_2^{(\mathcal{B})}$
$\mathcal{B}$ takes the following steps.
   (a) at time t_4, decrypts the content of $\hat{m}_2^{(\mathcal{C})}$, i.e., $m_2^{(\mathcal{C})} = \texttt{Dec}(\bar{k}_1, \hat{m}_2^{(\mathcal{C})})$.
   (b) at time t_5, checks the content of $m_2^{(\mathcal{C})}$. If $m_2^{(\mathcal{C})}$ is non-empty, i.e., $m_2^{(\mathcal{C})} = in_f$, it checks if the payee's detail in in_f has already been checked and the payment's amount does not exceed the customer's credit. If the checks pass, it runs the off-chain payment algorithm, $\texttt{pay}(in_f)$. In this case, it sets $m_2^{(\mathcal{B})}$="paid". Otherwise (i.e., if $m_2^{(\mathcal{C})} = \phi$ or neither checks pass), it sets $m_2^{(\mathcal{B})} = \phi$. It sends to $\mathcal{S}$ the encryption of $m_2^{(\mathcal{B})}$, i.e., $\hat{m}_2^{(\mathcal{B})} = \texttt{Enc}(\bar{k}_1, m_2^{(\mathcal{B})})$.

9. *Generating Complaint*. $\texttt{genComplaint}(\hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}, T, pk, aux_f) \rightarrow (\hat{z}, \hat{\ddot{\pi}})$
$\mathcal{C}$ takes the following steps.
   (a) decrypts $\hat{m}_1^{(\mathcal{B})}$ and $\hat{m}_2^{(\mathcal{B})}$; this results in $m_1^{(\mathcal{B})}$ and $m_2^{(\mathcal{B})}$ respectively. Depending on the content of the decrypted values, it sets its complaint's parameters $z := (z_1, z_2, z_3)$ as follows.
      - if $\mathcal{C}$ want to make one of the two below statements, it sets $z_1 =$"challenge message".
        (i) the "pass" message (in $m_1^{(\mathcal{B})}$) should have been a warning.
        (ii) $\mathcal{B}$ has not provided any message (i.e., neither pass nor warning) and if $\mathcal{B}$ provided a warning then the fraud would have been prevented.
      - if $\mathcal{C}$ wants to challenge the effectiveness of the warning (in $m_1^{(\mathcal{B})}$), it sets $z_2 = m||sig||pk_{\mathcal{G}}||$ "challenge warning", where m is an evidence, $sig \in aux_f$ is the evidence's certificate (obtained from the certificate generator $\mathcal{G}$), and $pk_{\mathcal{G}} \in pk$.

- if $\mathcal{C}$ wants to complain about the inconsistency of the payment (in $m_2^{(\mathcal{B})}$), then it sets z_3 = "challenge payment". Otherwise, it sets $z_3 = \phi$.

(b) at time t_6, sends to $\mathcal{S}$ the following values:
- the encryption of complaint z, i.e., $\hat{z} = \texttt{Enc}(\bar{k}_1, z)$.
- the encryption of $\ddot{\pi} := (\ddot{\pi}_1, \ddot{\pi}_2)$, i.e., $\hat{\ddot{\pi}} = \tilde{\texttt{Enc}}(pk_{\mathcal{D}}, \ddot{\pi})$. Note, $\ddot{\pi}$ contains the openings of the private statements' commitments (i.e., g_1, g_2), and is encrypted under each $\mathcal{D}_j$'s public key.

10. *Verifying Complaint*. $\texttt{verComplaint}(\hat{z}, \hat{\ddot{\pi}}, g, \hat{\boldsymbol{m}}, \hat{\boldsymbol{l}}, j, sk_{\mathcal{D}}, aux, pp) \rightarrow \hat{\boldsymbol{w}}_j$
Every Arbiter, $\mathcal{D}_j \in \{\mathcal{D}_1, ..., \mathcal{D}_n\}$, takes the following steps.
(a) at time t_7, decrypts $\hat{\ddot{\pi}}$, i.e., $\ddot{\pi} = \tilde{\texttt{Dec}}(sk_{\mathcal{D}}, \hat{\ddot{\pi}})$.
(b) checks the validity of $(\ddot{\pi}_1, \ddot{\pi}_2)$ in $\ddot{\pi}$ by locally running the SAP's verification, i.e., $\texttt{SAP.verify}(.)$, for each $\ddot{\pi}_i$. It returns s. If $s = 0$, it halts. If $s = 1$ for both $\ddot{\pi}_1$ and $\ddot{\pi}_2$, it proceeds to the next step.
(c) decrypts $\hat{\boldsymbol{m}} = [\hat{m}_1^{(\mathcal{C})}, \hat{m}_2^{(\mathcal{C})}, \hat{m}_1^{(\mathcal{B})}, \hat{m}_2^{(\mathcal{B})}]$ using $\texttt{Dec}(\bar{k}_1, .)$, where $\bar{k}_1 \in \ddot{\pi}_1$. Let $[m_1^{(\mathcal{C})}, m_2^{(\mathcal{C})}, m_1^{(\mathcal{B})}, m_2^{(\mathcal{B})}]$ be the result.
(d) checks whether $\mathcal{C}$ made an update request to its payee's list. To do so, it checks if $m_1^{(\mathcal{C})}$ is non-empty and (its encryption) was registered by $\mathcal{C}$ in $\mathcal{S}$. Also, it checks whether $\mathcal{C}$ made a payment request, by checking if $m_2^{(\mathcal{C})}$ is non-empty and (its encryption) was registered by $\mathcal{C}$ in $\mathcal{S}$ at time t_3. If either check fails, it halts.
(e) decrypts $\hat{z}$ and $\hat{\boldsymbol{l}}$ using $\texttt{Dec}(\bar{k}_1, .)$, where $\bar{k}_1 \in \ddot{\pi}_1$. Let $z := (z_1, z_2, z_3)$ and $\boldsymbol{l}$ be the result.
(f) sets its verdicts according to the content of $z := (z_1, z_2, z_3)$, as follows.
- if "challenge message" $\notin z_1$, it sets $w_{1,j} = 1$. Otherwise, it runs $\texttt{verStat}(add_{\mathcal{S}}, m_1^{(\mathcal{B})}, \boldsymbol{l}, \Delta, aux) \rightarrow w_{1,j}$, to determine if a warning (in $m_1^{(\mathcal{B})}$) should have been given (instead of the "pass" or no message).
- if "challenge warning" $\notin z_2$, it sets $w_{2,j} = w_{3,j} = 1$. Otherwise, it runs $\texttt{checkWarning}(add_{\mathcal{S}}, z_2, m_1^{(\mathcal{B})}, aux') \rightarrow (w_{2,j}, w_{3,j})$, to determine the effectiveness of the warning (in $m_1^{(\mathcal{B})}$).
- if "challenge payment" $\in z_3$, it checks whether the payment has been made. If the check passes, it sets $w_{4,j} = 0$. If the check fails, it sets $w_{4,j} = 1$. If "challenge payment" $\notin z_3$, it checks if "paid" is in $m_2^{(\mathcal{C})}$. If the check passes, it sets $w_{4,j} = 0$. Otherwise, it sets $w_{4,j} = 1$.

(g) encodes its verdicts $(w_{1,j}, w_{2,j}, w_{3,j}, w_{4,j})$ as follows.
i. locally maintains a counter, $o_{adr_{\mathcal{C}}}$, for each $\mathcal{C}$. It sets its initial value to 0.
ii. calls $\texttt{PVE}(.)$ to encode each verdict. In particular, it performs as follows. $\forall i, 1 \leq i \leq 4$:
- calls $\texttt{PVE}(\bar{k}_0, adr_{\mathcal{C}}, w_{i,j}, o_{adr_{\mathcal{C}}}, n, j) \rightarrow \bar{w}_{i,j}$
- $o_{adr_{\mathcal{C}}} = o_{adr_{\mathcal{C}}} + 1$.

By the end of this step, a vector $\bar{\boldsymbol{w}}_j$ of four encoded verdicts is computed, i.e., $\bar{\boldsymbol{w}}_j = [\bar{w}_{1,j}, ..., \bar{w}_{4,j}]$.
iii. uses $\bar{k}_2 \in \ddot{\pi}_2$ to further encode/encrypt $\texttt{PVE}(.)$'s outputs as follows. $\hat{\boldsymbol{w}}_j = \texttt{Enc}(\bar{k}_2, \bar{\boldsymbol{w}}_j)$.

(h) at time t_8, sends to $\mathcal{S}$ the encrypted vector, $\hat{\boldsymbol{w}}_j$.

11. *Resolving Dispute*. $\texttt{resDispute}(T_2, \hat{\boldsymbol{w}}, pp) \rightarrow \boldsymbol{v}$
$\mathcal{DR}$ takes the below steps at time t_9, when it is invoked by $\mathcal{C}$ or $\mathcal{S}$ which sends $\ddot{\pi}_2 \in T_2$ to it.
(a) checks the validity of $\ddot{\pi}_2$ by locally running the SAP's verification, i.e., $\texttt{SAP.verify}(.)$, that returns s. If $s = 0$, it halts. Otherwise, it proceeds to the next step.
(b) computes the final verdicts, as below.
i. uses $\bar{k}_2 \in \ddot{\pi}_2$ to decrypt the arbiters' encoded verdicts, as follows. $\forall j, 1 \leq j \leq n$:
$\bar{\boldsymbol{w}}_j = \texttt{Dec}(\bar{k}_2, \hat{\boldsymbol{w}}_j)$, where $\hat{\boldsymbol{w}}_j \in \hat{\boldsymbol{w}}$.
ii. constructs four vectors, $[\boldsymbol{u}_1, ..., \boldsymbol{u}_4]$, and sets each vector $\boldsymbol{u}_i$ as follows. $\forall i, 1 \leq i \leq 4$:
$\boldsymbol{u}_i = [\bar{w}_{i,1}, ..., \bar{w}_{i,n}]$, where $\bar{w}_{i,j} \in \bar{\boldsymbol{w}}_j$.
iii. calls $\texttt{FVD}(.)$ to extract each final verdict, as follows. $\forall i, 1 \leq i \leq 4$: calls $\texttt{FVD}(n, \boldsymbol{u}_i) \rightarrow v_i$.

(c) outputs $\boldsymbol{v} = [v_1, ..., v_4]$.

Customer $\mathcal{C}$ must be reimbursed if the final verdict is that (i) the "pass" message or missing message should have been a warning or (ii) the warning was ineffective and the provided evidence was not invalid, and (iii) the payment has been made. Stated formally, the following relation must hold:

$$\Big(\underbrace{(v_1 = 0)}_{\text{(i)}} \vee \underbrace{(v_2 = 0 \ \wedge \ v_3 = 0)}_{\text{(ii)}}\Big) \wedge \Big(\underbrace{v_4 = 0}_{\text{(iii)}}\Big),$$

where $v_i \in \boldsymbol{v}$.

Theorem 1. *The above PwDR scheme is secure, with regard to definition 7, if the digital signature is existentially unforgeable under chosen message attacks, the blockchain, SAP, and pseudorandom function* $\texttt{PRF}(.)$ *are secure, and the encryption schemes are semantically secure.*

5 Security Analysis of the PwDR Protocol

To prove the main theorem (i.e., Theorem 1), we show that the PwDR scheme satisfies all security properties defined in Section 3. We first prove that it meets security against a malicious victim.

Lemma 1. *If the digital signature is existentially unforgeable under chosen message attacks, the SAP and blockchain are secure, and the correctness of verdict encoding-decoding protocols (i.e., PVE and FVD) holds, then the PwDR scheme is secure against a malicious victim, with regard to Definition 4.*

Proof. First, we focus on event I : $\Big((m_1^{(\mathcal{B})} = warning) \wedge (\sum_{j=1}^{n} w_{1,j} < n)\Big)$ which considers the case where $\mathcal{B}$ has provided a warning message but $\mathcal{C}$ manages to make at least one of the arbiters set its verdict to 0, that ultimately results in $\sum_{j=1}^{n} w_{1,j} < n$. We argue that the adversary's success probability in this event is negligible in the security parameter. In particular, due to the security of SAP, $\mathcal{C}$ cannot convince an arbiter to accept a different decryption key, e.g., $k' \in \ddot{\pi}'$, that will be used to decrypt $\mathcal{B}$'s encrypted message $\hat{m}_1^{(\mathcal{B})}$, other than what was agreed between $\mathcal{C}$ and $\mathcal{B}$ in the initiation phase, i.e., $\bar{k}_1 \in \ddot{\pi}_1$. To be more precise, it cannot persuade an arbiter to accept a statement $\ddot{\pi}'$, where $\ddot{\pi}' \neq \ddot{\pi}_1$ except with a negligible probability, $\mu(\lambda)$. This ensures that honest $\mathcal{B}$'s original message is accessed by every arbiter with a high probability. Next, we consider event II : $\Big((\sum_{j=1}^{n} w_{1,j} = n) \wedge (v_1 = 0)\Big)$ that captures the case where all arbiters approved that the pass message was given correctly or the missing message could not prevent the APP fraud, but the final verdict that $\mathcal{DR}$ extracts does not imply that. We argue that the probability that this event occurs is negligible in the security parameter. Specifically, due to the security of the SAP, $\mathcal{C}$ cannot persuade (a) an arbiter to accept a different encryption key and (b) $\mathcal{DR}$ to accept a different decryption key other than what was agreed between $\mathcal{C}$ and $\mathcal{B}$ in the initiation phase. Specifically, it cannot persuade them to accept a statement $\ddot{\pi}'$, where $\ddot{\pi}' \neq \ddot{\pi}_2$ except with a negligible probability, $\mu(\lambda)$. Furthermore, as discussed in Section 4.2, due to the correctness of the verdict encoding-decoding protocols (i.e., PVE and FVD), the probability that multiple representations of verdict 1 cancel out each other is negligible too, $\frac{1}{2^{|p|}}$. Thus, the second event occurs only with a negligible probability.

Now, we move on to event III : $\Big((\texttt{checkWarning}(m_1^{(\mathcal{B})}) = 1) \wedge (\sum_{j=1}^{n} w_{2,j} < n)\Big)$. It captures the case where $\mathcal{B}$ has provided an effective warning message but $\mathcal{C}$ manages to make at least one of the arbiters set its verdict to 0, that ultimately results in $\sum_{j=1}^{n} w_{2,j} < n$. The same argument provided to event I is applicable to this even too. Briefly, due to the security of SAP, $\mathcal{C}$ cannot persuade an arbiter to accept a different decryption key other than what was agreed between $\mathcal{C}$ and $\mathcal{B}$ in the initiation phase; therefore, all arbiters will receive the original message of $\mathcal{B}$, except a negligible probability, $\mu(\lambda)$. Now, we consider event IV : $\Big((\sum_{j=1}^{n} w_{2,j} = n) \wedge (v_2 = 0)\Big)$, which captures the case where all arbiters approved that the warning message was effective but the final verdict that $\mathcal{DR}$ extracts does not indicate that. The security argument of event II applies to this event as well. In short, due to the security of the SAP, $\mathcal{C}$ cannot persuade an arbiter to accept a different encryption key, and cannot convince $\mathcal{DR}$ to accept a different decryption key other than what was initially agreed between $\mathcal{C}$ and $\mathcal{B}$, except a negligible probability, $\mu(\lambda)$. Furthermore, due to

the correctness of PVE and FVD protocols, the probability that multiple representations of verdict 1 cancel out each other is negligible, $\frac{1}{2^{|p|}}$. Thus, the probability that this event takes place is negligible. Now, we analyse event V : $\Big(u \notin Q \wedge \texttt{Sig.ver}(pk, u, sig) = 1\Big)$. This even captures the case where the malicious victim comes up with a valid signature/certificate on a message that has never been queried to the signing oracle. Nevertheless, due to the existential unforgeability of the digital signature scheme, the probability that such an event occurs is negligible, $\mu(\lambda)$. Next, we focus on event VI : $\Big((\sum_{j=1}^{n} w_{3,j} < n) \wedge (v_3 = 0)\Big)$ that considers the case where a subset of arbiters indicated that the signature (in $\mathcal{C}$'s complaint) is invalid, but the final verdict that $\mathcal{DR}$ extracts implies that all arbiters approved it. This means the adversary has managed to switch the verdicts of those arbiters which voted 0 to 1. However, the probability that this even occurs is negligible as well. Because, due to the SAP's security, $\mathcal{C}$ cannot convince an arbiter and $\mathcal{DR}$ to accept different encryption and decryption keys other than what was initially agreed between $\mathcal{C}$ and $\mathcal{B}$, except a negligible probability, $\mu(\lambda)$. Therefore, with a negligible probability the adversary can switch a verdict for 0 to the verdict for 1.

Moreover, a malicious $\mathcal{C}$ cannot frame an honest $\mathcal{B}$ for providing an invalid message by manipulating the smart contract's content, e.g., by replacing an effective warning with an ineffective one in $\mathcal{S}$, or excluding a warning from $\mathcal{S}$. In particular, to do that, it has to either forge the honest party's signature, so it can send an invalid message on its behalf, or fork the blockchain so the chain comprising a valid message is discarded. In the former case, the adversary's probability of success is negligible as long as the signature is secure. The adversary has the same success probability in the latter case, because it has to generate a long enough chain that excludes the valid message which has a negligible success probability, under the assumption that the hash power of the adversary is lower than those of honest miners and due to the blockchain's liveness property an honestly generated transaction will eventually appear on an honest miner's chain [6]. □

Now, we first present a lemma formally stating that the PwDR scheme is secure against a malicious bank and then prove this lemma.

Lemma 2. *If the SAP and blockchain are secure, and the correctness of verdict encoding-decoding protocols (i.e., PVE and FVD) holds, then the PwDR scheme is secure against a malicious bank, with regard to Definition 5.*

Proof. We first focus on event I : $\Big((\sum_{j=1}^{n} w_{1,j} < n) \wedge (v_1 = 1)\Big)$ which captures the case where

□

5.1 Extensions

Micro-enterprise or Charity Customer. There is a clause in the CRM code that allows a bank to refuse reimbursing a victim of an APP scam, if (i) the victim is an organisation (i.e., Micro-enterprise or Charity), (ii) the organisation has internal payment procedures that are effective in preventing the APP scam, and (iii) the victim has not followed those procedures. Below, we outline how the PwDR protocol can be extended to capture these requirements. The modified PwDR allows a customer to prove it has followed those requirements (or to prove there were not such procedures), which ultimately benefits an honest customer during the dispute resolution phase (i.e., phase 11). We present the extension in two phases. In phase I, we provide a subprotocol that determines whether the customer has met the above requirements. In phase II, we show how the subprotocol can be integrated into the PwDR protocol.
Phase I. In this phase, we provide an overview of a subprotocol that determines if the customer has met the above requirements. At a high level, this subprotocol works as follows.

(A) $\mathcal{C}$ sends the organisation's internal procedure specification and the specification's certificate to $\mathcal{S}$. Moreover, it sends to $\mathcal{S}$ a proof, p, asserting that it has followed the above procedure.
(B) Every $\mathcal{D}_i$ takes the following steps:

(a) checks the certificate by running procedure $\texttt{Sig.ver}(.) \rightarrow h_i$. It sends h_i to $\mathcal{S}$. It waits until all arbiters provide their inputs. Then, it locally runs $\texttt{f.verdict}(h_1, ..., h_n)$ to determine if the certificate has been approved (by the majority of arbiters). It only proceeds to the next step if $\texttt{f.verdict}(.)$ returns 1.

(b) checks whether the procedure could prevent the APP scam. If the check passes, then it sends $v'_i = 1$ to $\mathcal{S}$; otherwise, it sends $v'_i = 0$ to $\mathcal{S}$. Again, it waits until all arbiters provide their inputs. Next, it locally runs $\texttt{f.verdict}(v'_1, ..., v'_n)$. It only proceeds to the next step if $\texttt{f.verdict}(.)$ returns 1.

(c) verifies proof p. If the check passes, then it sends $v''_i = 1$ to $\mathcal{S}$. Otherwise, it sends $v'_i = 0$ to $\mathcal{S}$.

(C) $\mathcal{S}$ computes and stores the following values:

- $h = \texttt{f.verdict}(h_1, ..., h_n)$.
- if $h = 1$, then $v' = \texttt{f.verdict}(v'_1, ..., v'_n)$.
- if $v' = 1$, then $v'' = \texttt{f.verdict}(v''_1, ..., v''_n)$.
- sets $g = 1$, if one of the following two conditions holds:
  * $h = 0$.
  * $h \wedge v'' = 1$.

  otherwise (if neither condition holds), sets $g = 0$.

Phase II. Now describe how the above subprotocol's phases can be integrated into the PwDR protocol. First, phases (A) and (B), of the subprotocol, are added to steps 9a and 10f, of the PwDR protocol, respectively. Second, the checks in phase **??** of the PwDR protocol need to also ensure that $g = 1$, by taking the steps of phase (C) in the subprotocol.

Note that in step (B)b of the subprotocol, it is assumed that each arbiter uses a well-defined process to evaluate whether the customer's internal payment procedures could have prevented the APP scam. Nevertheless, effective payment procedures that prevent the APP fraud have not been appropriately defined by the CRM code. This leads to a manual and inefficient evaluation process. Therefore, one may ask:

Q1: which measures exactly should be included in the internal payment procedures of an organisation to prevent APP scams?

Given an explicit list of payment procedures that prevent APP scams, we could make the above evaluation procedure autonomous and transparent. In particular, such a list could be encoded into the smart contract $\mathcal{S}$ which receives inputs from the customer (and possibly arbiters) and check if those internal payment procedures were effective in preventing the APP scam.

Security against exploitative victim. Having in place a transparent deterministic procedure (e.g., the PwDR protocol) for evaluating victims' requests for reimbursement could potentially create opportunities for exploitations. In particular, an honest victim of the APP fraud who had been reimbursed in the past due to the payment system's vulnerability, e.g., ineffective payment, may be tempted to exploit the same known vulnerability multiple times. Below, we outline two cases in which the malicious victim may exploit a known vulnerability and describe how they can be dealt with.

- Case 1: A genuine victim, who previously had been reimbursed, colludes with a certain payee. In this case, the malicious victim uses the previous working strategy (e.g., the same certificates and complaints) to declare that it has been a victim of the APP fraud but this time it exploits the weakness (e.g., ignores the ineffective warning). Note, the colluding payee acts exactly the same way as a real scammer does in the APP scam, e.g., after receiving the payment, it transfers the money to another account abroad.
- Case 2: A genuine victim shares (or sells) its knowledge of the vulnerability to a malicious customer who colludes with a payee to claim it has been a victim of the APP scam. Therefore, this case is similar to Case 1 with the main difference that the claimant's identity changes each time a claim is made.

Ideally, the bank should patch the vulnerability as soon as the first incident occurs; however, this may not be the case in the real world as the cost of improving a (payment) system to deal with such an incident may far exceed the bank's monetary loss in that incident. Currently, neither the CRM code nor the "Payment Services Regulations" [13] explicitly offer any solution for the above cases. To address the issue in Case 1, we propose the following remedy. First, we introduce two (set of) parameters; namely, threshold and counters. We let bank $\mathcal{B}$ define the value of the threshold in the smart contract, $\mathcal{S}$. Also, we require the protocol to keep track of the number of times the same customer is reimbursed for the same complaint, e.g., ineffective warning. Now we outline how these parameters can assist the protocol to rectify the issue. Each time the protocol receives a customer's complaint, it checks whether the total number of times the same customer is reimbursed for that specific complaint exceeds the predefined threshold; if so, the protocol discards that complaint. Otherwise, it proceeds as before. More specifically, we amend the PwDR protocol as follows.

- $\mathcal{S}$ maintains a counter vector, $\boldsymbol{q}^{(\mathcal{C})} = [(q_e, e), (q_x, x), (q_y, y)]$, for each customer $\mathcal{C}$, where q_e, q_x, and q_y are initially set to 0 while e, x and y are the types of complaint, as they were defined in the PwDR scheme.
- $\mathcal{B}$ defines a fixed threshold t in $\mathcal{S}$, where $t > 1$.
- In Phase 11, $\mathcal{S}$ increments counter q_j by 1, each time $\mathcal{C}$ is reimbursed for complaint j, where $j \in \{e, x, y\}$.
- Algorithm `verComplaint`(.) takes extra parameters $\boldsymbol{q}^{(\mathcal{C})}$ and t as inputs.
- In Phase 10, when a committee member, $\mathcal{D}_i$, wants to examine $\mathcal{C}$'s complaint, it reads the content of $\boldsymbol{q}^{(\mathcal{C})}$ and t from $\mathcal{S}$ and passes them to `verComplaint`(.), which first checks whether the counter in $\boldsymbol{q}^{(\mathcal{C})}$ exceeds t for the complaint that $\mathcal{C}$ is making. If the check passes, it sets $d_i = v_i = \bar{v}_i = w_i = \phi$ and outputs $(d_i, v_i, \bar{v}_i, w_i)$ without requiring $\mathcal{D}_i$ to process $\mathcal{C}$'s claim. Otherwise, $\mathcal{D}_i$ checks the claim as before.

Now, we move on to the issue in Case 2 which is harder to identify than the one in Case 1. Because, in the former case, the identity of a malicious customer which wrongly claims that it has fallen victim to the APP fraud changes continuously and is hardly distinguishable from a genuine victim without the assistance of extra information (e.g., phone records, emails) that is not trivial to attain without external intervention, e.g., a subpoena. To address this issue, we propose the following mitigation. Briefly, we require $\mathcal{B}$ to enhance its system and rectify the issue when the number of complaints related to the same issue exceeds a global threshold. Specifically, $\mathcal{B}$ defines in $\mathcal{S}$ a *global* threshold gt and global counter vector $\boldsymbol{g} = [(g_e, e), (g_x, x), (g_y, y)]$, where the vector's elements are defined the same way as those are defined in $\boldsymbol{q}$, with the difference that they are global and are not for a specific customer. Each time a customer is reimbursed for a reason the related counter in $\boldsymbol{g}$ is incremented by 1. When a counter exceeds the value of gt, $\mathcal{S}$ notifies $\mathcal{B}$ which patches the issue and then sets the related counter to 0. Note that the value of gt can be set such that when the monetary loss exceeds the upgrade costs, then $\mathcal{B}$ upgrades the system.

6 Discussion

6.1 Effective Warning

In the protocol, in step 7, to which value $\mathcal{C}$ sets its message $m_2^{(\mathcal{C})}$ is not totally deterministic and depends on various (external) factors, e.g., warning effectiveness, human factors. It is very likely that if $m_1^{(\mathcal{B})} =$ "pass", then $\mathcal{C}$ acts deterministically, by asking $\mathcal{B}$ to make the payment, i.e., $\mathcal{C}$ sets $m_2^{(\mathcal{C})} = in_p$. However, to which value $\mathcal{C}$ sets $m_2^{(\mathcal{C})}$ when (a) $m_1^{(\mathcal{B})} =$ warning or (b) no message is provided by $\mathcal{B}$ at time T_2, depends on many factors. For instance, if the warning is not effective and does not concern $\mathcal{C}$, then $\mathcal{C}$ still would set $m_2^{(\mathcal{C})} = in_p$. Similarly, if $\mathcal{B}$ does not send $m_1^{(\mathcal{B})}$, it is still possible that $\mathcal{C}$ sets $m_2^{(\mathcal{C})} = in_p$, for instance when it is distracted and does not pay attention to the absence of the message that was supposed to be provided by $\mathcal{B}$ on time. On the other hand, if the warning is effective, or $\mathcal{C}$ in general is highly sensitive to warnings and the absence of $\mathcal{B}$'s message, then it does not make any payment, i.e., it sets $m_2^{(\mathcal{C})} = \phi$. Despite the above challenges, the proposed protocol ensures that $\mathcal{C}$ will be identified as the party who should be reimbursed, if it acts according to the protocol, makes a payment but the warning was ineffective or no message was provided by $\mathcal{B}$ in step 7. Therefore, the following questions would follow:

Q2: what percentage of customers after encountering the warning do still proceed to make a payment (i.e., set $m_2^{(\mathcal{C})} = in_p$)?

Q3: how to make the above rate negligibly small?

There exists a comprehensive research line in determining the effectiveness of warnings, e.g., in [8,2,5]. These traditional research line studies which factors make warnings effective and how a warning recipient is attracted to and follows the warning message. However, in the context of APP scams, there is a vital unique factor that can directly influence a warning effectiveness; the factor is *the ability of the scammer to interact directly with the victim* (or warning recipient). This lets a scammer to actively try to negate the effectiveness of banks' warnings and persuade the warning recipient to ignore the warning (and make payment). Such a factor was not (needed to be) taken into account in the traditional study of warnings. Therefore, one may ask:

Q4: to what extent can a scammer negate a warning effectiveness in the onctext of APP scams?

One way to answer the above question is to study each individual bank's statistics and find out how successful they have been in combatting the APP scams, as each bank designs its payment system and warnings independent of other banks. Thus, it would be interesting to find out:

Q5: which bank does have the lowest rate of APP scams and how did its warnings contribute to the low rate?

Furthermore, in step 10f, it is implicitly assumed that in order for each arbiter, $\mathcal{D}_i$, to judge the effectiveness of the warning and reach a verdict, it has access to the payment system and it can interact with $\mathcal{C}$ offline, e.g., to obtain further evidence from it. Nevertheless, this process can be time consuming and a verdict may not be released in the real time once $\mathcal{C}$ sends its complaint to $\mathcal{S}$. Thus, it is natural to ask:

Q6: to what extent can the role of the arbiters be automated and accurately played by a computer program?

6.2 Ensuring the Payment is for Genuine Goods and Services

There are a set of terms in the CRM code that states:
"In all the circumstances at the time of the payment, in particular the characteristics of the Customer and the complexity and sophistication of the APP scam, the Customer made the payment without a reasonable basis for believing that:

(i) *the payee was the person the Customer was expecting to pay;*
(ii) *the payment was for genuine goods or services; and/or*
(iii) *the person or business with whom they transacted was legitimate".*

We argue that clause (ii) plays a minor role in preventing a APP scam, as it is effective only when the seller wrongly claims it is in possession of a certain goods or can deliver certain services. It is ineffective when a customer ensures goods or services it wants to receive are indeed genuine but the seller avoids delivering the goods/services. In this case, a seller (regardless of whether it is legitimate or not) may prove to the customer that the goods and services are genuine and belong to it (e.g., by sending a copy of related genuine document), but it still avoids delivering them once it receives the money. To capture the above issue as well, the above clause should be modified as follows:

"... the delivery of genuine goods or services are guaranteed...".

There are at least two ways to have the above guarantee: (a) involving a reputable trusted third-party intermediary which can compensate the customer if the seller misbehaves, e.g., eBay, amazon, or (b) using a secure "contingent service payment" scheme (e.g., in [3]) that supports the fair exchange of digital goods or services and money without the involvement of the above trusted third-party intermediary. In this case, there

would be no need for the customer to ensure if the seller is legitimate, because it pays only if genuine goods or services are delivered. The existing fair exchange schemes allow a seller and buyer to trade with each other *outside of the bank payment system* by using a blockchain. However, these schemes can be embedded into the bank payment system such that the bank (or a group of banks) maintains the blockchain and converts internally a customer's fiat currency to a digital one when the customer wants to trade in a fair manner with a party whose authenticity cannot be verified.

6.3 Ensuring the Legitimacy and Authenticity of Payee

We highlight that even though clauses (i) and (iii), presented in Section 6.2, can prevent a customer from falling to an APP scam, only in certain cases they would benefit a victim of an APP fraud during the process of allocating liability or dispute resolution. In particular, if the victim declares that (a) it has used a secure authentication mechanism (e.g., digital signature) to ensure the legitimacy and authenticity of the payee or (b) it has not performed any authentication, then it would not receive the reimbursement. In the former case, it has transferred the money to the party it knows. However, this is not an APP scam, according to the scam's general definition. In the latter case, the firm can avoid reimbursing it, according to the CRM code (as the customer has not performed its part). The only case in which the customer might be reimbursed is when it (a) uses an insecure means for authentication (e.g., phone call) and (b) can prove it has done such a check. For instance, in the case where the victim has been asked by a scammer to make a phone call to its bank while the scammer answers the phone call that the victim makes later. In this case, the victim needs to prove that it has made a phone call to the bank, which may not be always possible (e.g., if the scammer uses a method to reply to the call before the real connection between the customer and bank is made). In this case, in order for the customer to be reimbursed it needs to provide an evidence while there is no evidence left behind by the scammer. Thus, it is natural to ask:

Q7: how can the victim prove it has been a victim of an APP scam, in the above case?

6.4 Lack of Gross Negligence Definition

One of the conditions in the CRM code that allows a bank to avoid reimbursing the customer is clause R2(1)(e) which states:

"The Customer has been grossly negligent. For the avoidance of doubt the provisions of R2(1)(a)-(d) should not be taken to define gross negligence in this context."

Nevertheless, neither the CRM code nor the Payment Services Regulations explicitly define under which circumstances the customer is considered "grossly negligent" in the context of the APP scam. In particular, in the CRM code, the only terms that discuss customer's misbehaviour are the provisions of R2(1)(a)-(d); however, as stated above, they should be excluded from the definition of the term gross negligence. On the other hand, in the Payment Services Regulations, this term is used three times, i.e., twice in regulation 75 and once in regulation 77. But in all three cases it is used for frauds related to *unauthorised payments* which are different types of frauds from the APP scams. Therefore, even the Payment Services Regulations does not define the term in the context of the APP scam.

When an accurate definition of the term is in place, its conditions can be encoded into the smart contract of the PwDR protocol. This allows the PwDR protocol to transparently resolve disputes between the bank and customer, if the bank claims that the customer has been grossly negligent.

References

1. Abadi, A., Murdoch, S.J., Zacharias, T.: Recurring contingent payment for proofs of retrievability. Cryptology ePrint Archive, Report 2021/1145 (2021), `https://ia.cr/2021/1145`

2. Brinton Anderson, B., Vance, A., Kirwan, C.B., Eargle, D., Jenkins, J.L.: How users perceive and respond to security messages: a neurois research agenda and empirical study. European Journal of Information Systems 25(4), 364–390 (2016)
3. Campanelli, M., Gennaro, R., Goldfeder, S., Nizzardo, L.: Zero-knowledge contingent payments revisited: Attacks and payments for services. In: CCS'17
4. Confirmation of Payee Team: Confirmation of payee: Response to consultation cp20/1 and decision on varying specific direction 10 (2020), https://www.psr.org.uk/media/qrbc03jm/psr-ps20-1-variation-of-specific-direction-10-february-2020.pdf
5. Felt, A.P., Reeder, R.W., Almuhimedi, H., Consolvo, S.: Experimenting at scale with google chrome's ssl warning. In: Proceedings of the SIGCHI conference on human factors in computing systems (2014)
6. Garay, J.A., Kiayias, A., Leonardos, N.: The bitcoin backbone protocol: Analysis and applications. In: EUROCRYPT (2015)
7. Katz, J., Lindell, Y.: Introduction to Modern Cryptography, Second Edition. CRC Press (2014), https://www.crcpress.com/Introduction-to-Modern-Cryptography-Second-Edition/Katz-Lindell/p/book/9781466570269
8. Laughery, K.R., Wogalter, M.S.: Designing effective warnings. Reviews of human factors and ergonomics (2006)
9. Lending Standards Board: Contingent reimbursement model code for authorised push payment scams (2021), https://www.lendingstandardsboard.org.uk/wp-content/uploads/2021/04/CRM-Code-LSB-Final-April-2021.pdf
10. Nakamoto, S.: Bitcoin: A peer-to-peer electronic cash system. Tech. rep. (2019)
11. Pedersen, T.P.: Non-interactive and information-theoretic secure verifiable secret sharing. In: CRYPTO (1991)
12. Schneier, B.: Applied cryptography - protocols, algorithms, and source code in C, 2nd Edition. Wiley (1996)
13. The Financial Conduct Authority: The payment services regulations (2017), https://www.legislation.gov.uk/uksi/2017/752/contents/made
14. Wood, G., et al.: Ethereum: A secure decentralised generalised transaction ledger. Ethereum project yellow paper (2014)